TUM

TUM SCHOOL OF COMPUTATION, INFORMATION AND TECHNOLOGY

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

eBPF-Assisted Relays for Multimedia Streaming

Daniel Alexander Antonius Pfeifer

TUM

TUM SCHOOL OF COMPUTATION, INFORMATION AND TECHNOLOGY

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

eBPF-Assisted Relays for Multimedia Streaming

eBPF-Unterstützung für Multimedia-Streaming-Netzknoten

Author:	Daniel Alexander Antonius Pfeifer
Supervisor:	Prof. Dr.-Ing. Jörg Ott
Advisor:	Mathis Engelbart, M.Sc.
Submission Date:	15.08.2024

I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

Munich, 15.08.2024   Daniel Alexander Antonius Pfeifer

Acknowledgments

Abstract

In this thesis we propose a new relay setup for multimedia streaming that allows for avoidance of userspace processing by utilizing BPF programs in the Linux kernel. In a sample implementation we demonstrate the feasibility of this approach by designing a relay that is capable of forwarding packets between a server-side and a client-side QUIC connection while still being able to do adaptive bitrate streaming based on client congestion.

We show that this approach saves processing time and reduces latency compared to userspace processing with the relay still adhering to specifications of the QUIC standard and the 'Media over QUIC' (MoQ) draft. One limitation that is not addressed in depth in this thesis is the need for a de- and encryption hardware offload onto a SmartNIC to allow the BPF program to access the packet payload without any restrictions. Since the QUIC standard is still fairly new we are confident that a solution for a potential hardware offload will be found in future research.

Contents

1 Introduction

The fact that online streaming tends to be slower than cable TV is likely something most people have already experienced first-hand. Live sports events, music shows or news broadcasts arrive an order of seconds later when streamed compared to using traditional cable connections. Now despite this delay not necessarily being a deal-breaker for most people, designing networks that tighten the gap between cable and streaming is still a worthwhile endeavor. With very optimized and fast networks already in place, we are at a point where providing such a faster information delivery is highly non-trivial. To do that we have even gone as far as developing completely new standards, such as QUIC. Those new standards aim to improve the shortcomings of some of the most fundamental protocols of the internet, one of them being TCP, which have been around more than 40 years.

Besides introducing new protocols, one could also look at existing setups and figure out how to trade some generality for a smaller delay when handling data. The ISO/OSI model, which is a foundational concept in networking, provides "a common basis for the coordination of standards development for the purpose of systems interconnection" [Sta94]. As one can imagine, such a "common basis", even though convenient for large scale systems, can cause unnecessary overhead. In some cases additional speed-ups can be achieved by using more application-specific approaches. This thesis will consider one of such cases and explore the possibilities of avoiding or delaying certain processing steps of the ISO/OSI model in order to increase the overall speed of a network relay.

1.1 Research Question

As already mentioned above the usage of application specific approaches in networking allows for a reduction in latency. In this thesis we will consider a media streaming scenario that runs on top of QUIC by using the "Media over QUIC" (MoQ) transport protocol [Int24a]. The central question we will try to answer in this thesis will then be:

How can we improve the performance of a relay in a media streaming scenario by using eBPF technology?

with more specific sub-questions being:

1. *How can we avoid the need to direct a packet through userspace?*

2. *How to handle the fact that packets are heavily encrypted?*

3. *What communication between userspace and eBPF program is necessary to stay coherent with potential state?*

4. *How can our approach be generalized to other protocols?*

By using eBPF technology together with kernel hook points provided by the Linux-Kernel, we will try to find a setup that improves relay performance using eBPF programs that handle basic relay capabilities, such as packet forwarding and congestion control. Since the QUIC protocol is designed to handle a large portion of its workload in userspace we look into possibilities of delaying any userspace processing until **after** the packet has been forwarded to the client. This way the raw delay that the packet experiences from the initial media server to the client could be reduced. However, since QUIC is a connection oriented protocol, we need to make sure that the QUIC connection state stays coherent despite the additional processing steps done by the eBPF program. We will investigate which additional processing steps are needed in our case, how they compare to challenges when expanding our approach to other protocols and how they can be implemented in an eBPF program.

1.2 Scope

The main improvement this thesis aims to achieve is shortening the critical path a packet takes from a media server to a client. This will be done by avoiding the immediate need of a packet traversal up the network stack to the application layer. Instead, any communication with the application layer will happen in a delayed fashion (after the packet was sent) by utilizing eBPF-Maps for storing any necessary (meta-) information. The main reason this communication between userspace and eBPF-program is required lies in the fact that relays in MoQ are an application layer concept. That means the QUIC connections to from relay to server and from relay to client will be different and the packets that have been eBPF-forwarded to egress directly will need changes in their header data in order to match the state of the outgoing client connection.

This approach is highly dependent on the used standards and protocols. This thesis will operate on top of the QUIC protocol [Int24b] and the "Media over QUIC" (MoQ)

transport protocol [Int24a]. For the application layer the quic-go library [See24] will provide the implementation and any additional (non-eBPF) program will also be written in Go. Since the setup is dependent on retreiving data from eBPF-Maps the QUIC library providing the implementation will need some adaptations. We will mainly introduce simple function pointer style additions that allow the adapted library to be run both with and without the eBPF setup. The developer of the relay will then also have more freedom to setup the eBPF part of the relay as they see fit since the Go code that will interact with eBPF parts will also have to be provided by said developer.

Additionally we will run a performance analysis on our implementation of the relay to confirm the potential this approach has. These performance tests will look at the raw delay speedup as well as the impact on CPU utilization this setup has. All the tests will be done in a lab-like environment to isolate the performance changes as best as possible from any outside noise. The payloads used will only contain dummy data since our approach does not interfer with payload contents and there is no need for creating and using real media stream data.

Despite our approach only considering QUIC and MoQ, we will argue that the general idea of our setup will be independent of any of these protocols and can be changed to fit ones needs.

With this we will provide answers to the research questions regarding packet-redirection, communication between userspace and eBPF as well as setup-generalization. Regarding the question on how to handle the encryption of the packets, we will not focus on this since we did not find a suitable hardware offload that would have allow for en- and decryption after and before the used eBPF hook points respectively. Instead we will emulate this behavior by turning off the encryption in the QUIC library itself which will provide a similar result.

1.3 Structure of this Thesis

In chapter 2 this thesis will provide some overview of used technologies and related ideas. Section 2.1 will give an introduction to the QUIC protocol and its main features and section 2.2 will provide an overview of eBPF technology together with features related to our approach. Section 2.3 will introduce the 'Media over QUIC' (MoQ) transport protocol which will be used for our application level relay setup. After that section 2.4 will explain the ideas and challenges of adaptive bitrate streaming while section 2.5 will mention some work related to the aforementioned topics. What will follow in chapter 3 is a detailed description of the setup that allowed us to improve relay performance. We will look at the adaptations to the used QUIC library in section 3.1 as well as our eBPF setup in 3.2. Besides those two we will also look at some more specific

details and challenges in the subsequent sections. In chapter 4 we will then provide a basic performance analysis of our setup to show current improvements and limitations. Finally we will conclude with a summary together with some ideas for future work in this field in chapters 5 and 6.

1.4 Citation Examples

Citation [Int24b]. Citation [Sta94]. Citation [Int24a]. Citation [Rod24]. Citation [W3T24]. Citation [JC20]. Citation [Lie18]. Citation [Yan+20]. Citation [Lan+17]. Citation [PVV21]. Citation [KWF03]. Citation [DB19]. Citation [Tyu22]. Citation [See24]. Citation [Pfe24a]. Citation [Pfe24b]. Citation [PE24]. Citation [Pfe24c]. Citation [Fou].

2 Background and Related Work

2.1 QUIC

The Transmission Control Protocol (TCP) has been used as the backbone of the internet for more than 40 years. It has been designed to be reliable and to provide a connection-oriented way of transmitting data, but the modern environment of the internet with its need for increasing throughput makes it hard for TCP to keep up. Limitations in the design and resulting issues like head-of-line blocking have raised demand for a newly designed protocol that can keep up with the modern internet. The 'Quick UDP Internet Connections' protocol, short QUIC protocol, is a transport layer protocol built on top of UDP that is designed to be reliable, cryptographically secure and more performant than TCP. It was intended to be the successor of TCP and it has its origins at Google before its standardization by the IETF began in 2016. QUIC, partly because it operates both in user- and kernel-space, has been designed to allow for a more rapid deployment cycle than TCP. Similar to TCP it is a connection-based protocol that uses TLS for encryption [Lan+17]. Already back in 2018, QUIC was the default protocol for the Google Chrome browser which, at the time, made up 60% of the web browser market [Lie18]. A little over two years later, Facebook, now Meta, was using QUIC for more than 75% of their internet traffic which led to improvements regarding request errors, tail latency and header size [JC20]. As of May 2024, QUIC already made up 8.0% of all internet traffic with support from pretty much every major browser [W3T24; Rod24]. With big players like Google, Meta, Microsoft or YouTube putting emphasis on using QUIC to improve their services, this number is likely to increase even further.

2.1.1 Connections and Streams

Since QUIC is a connection-based protocol, some initial overhead to establish a connection is needed. However, the design incorporates some features that aim for a more efficient way of establishing connections, e.g. by using 0-RTT (zero round-trip-time) handshakes. Latency improvements like the 0-RTT handshake however come at the cost of security, since that opens the door for replay attacks. Another part where QUIC tries to optimize connection management is the use of streams. Streams are designed to be lightweight and can be opened without the need of a handshake. This goes as far as

a single packet being able to open a new stream, transferring stream data as well as the closing the stream agian. This allows for new techniques to improve data transmission and will also be part of the fast-relay setup in this thesis. Aside from streams, apparent since QUIC is based on UDP, it is also possible to send data via unreliable datagrams. This further improves versatility of the protocol and allows for new ways of optimizing data transmission.

2.1.2 quic-go and moqtransport

The implementation of the proposed fast-relay setup will be based on a modified version of the quic-go library, which provides a pure Go implementation of the QUIC protocol as specified in the standards RFC-9000, RFC-9221 as well as some others which are not important for out usecase. Besides that the fast-relay implementation will also make use of the moqtransport library. This library brings the 'Media over QUIC' (MoQ) protocol to Go and will be used as a media transport protocol when looking at the impact of fast-relays on adaptive real-time video streaming. The MoQ protocol is being standardized by the IETF since July 2023 and has yet to be finalized.

2.1.3 QUIC and Fast-Relays

The QUIC protocol will be a fundamental part of the fast-relay setup in this thesis, yet the ideas used to make relays faster is not limited to QUIC and can be extended to other protocols as well. QUIC is chosen as an example protocl due to its increasing popularity which offers big potential in early adoption and deployment of fast-relays. Besides that, the existing implementations of QUIC related standards provide a good starting point for an implementation, despite the difficulties that the heavy encryption of QUIC brings with it. To mitigate missing technologies, mainly for offloading QUIC decryption and encryption onto hardware, the existing protocol libraries can also be modified easily to simulate any needed behavior.

2.2 eBPF

In 1992 a technology called 'Berkeley Packet Filter' (BPF) was introduced into the Unix kernel. By using BPF it is possible to attach a small BPF-program to some pre-defined hook points in the network stack of the kernel and filter packets there in a stateless manner. This provided more efficiency since the packets did not need to be copied into userspace anymore but could directly be processed in the kernel. One downside to such an approach however is that BPF-programs are limited by the so-called 'BPF-verifier' which needs to check every BPF-program for safety e.g. to avoid infinite loops or access

to invalid memory from withing kernel space. Today, the initial technology of BPF has evolved into 'extended BPF' (eBPF) and allows for more versatile use cases.

2.2.1 eBPF Hook Points

The Linux kernel offers several hook points where eBPF-programs can be attached to. Two main ones are one to attach eBPF-programs to the Traffic Control (TC) subsystem and another one to attach them to the eXpress Data Path (XDP) subsystem. The XDP hook, which is directly located in the NIC-driver, lies lower in the network stack than the TC-hook, which is located in the link-layer. Despite being higher up in the network stack, the TC-hook has the big advantage that it offers ingress and egress processing while the XDP-hook is available for ingress processing only. This makes the XDP-hook suboptimal for the implementation of fast-relays since they heavily rely on processing packets at egress, after those were redirected from ingress. Figure Figure 2.1 illustrates again the relative positions of the TC and XDP hook points in the network stack.

2.2.2 Traffic Control Queuing Disciplines

The Linux Traffic Control Subsystem uses Queuing Disciplines (qdiscs) to define how packets are handled. TODO

2.2.3 eBPF Verifier

Since eBPF programs essentially allow the user to run arbitrary code in the kernel, it is crucial to ensure that this code is safe. This is achieved by the eBPF verifier, which checks the code for, and disallows, insecure sections. For example this includes not allowing unbounded loops or those with conditional exits. The verifier does find such insecurities by simulating possible program paths as well as looking at a graph representation of the program [Fou].

2.2.4 Important eBPF Concepts

One of the most important concepts in eBPF is the 'eBPF-map'. Such a map boils down to a section in memory that is reserved for the eBPF-program and which can be used as a key-value store for arbitrary data. This part of memory can then also be accessed from userspace and thus provide the main way of communication between the eBPF-program and our application. When we define an eBPF-map we can choose between different types as well as configure size, key-type, value-type and the way the map is stored.

```
1 // TODO: add example of eBPF map
```
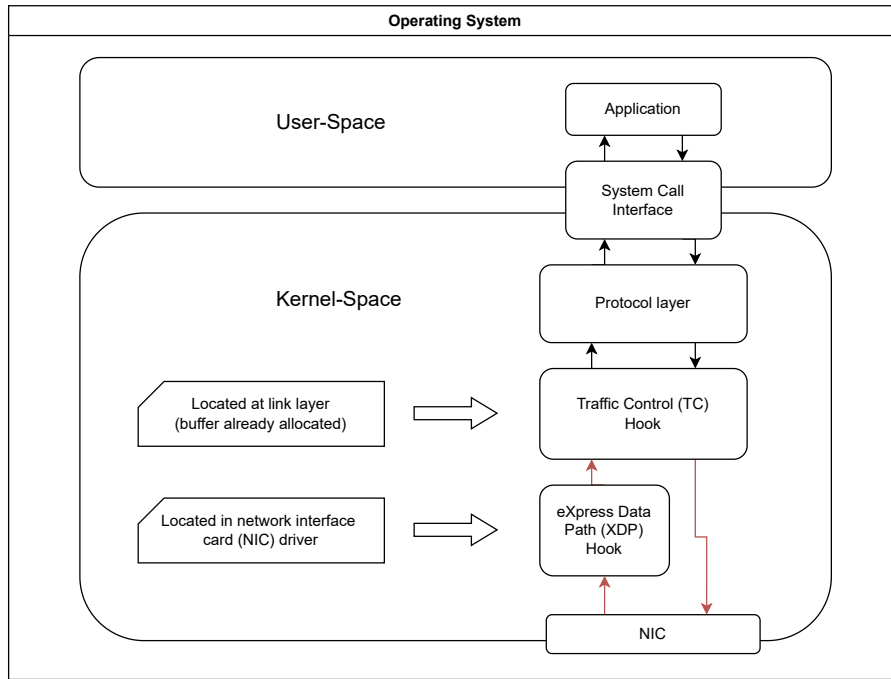
Figure 2.1: Abstracted view of Traffic Control (TC) and eXpress Data Path (XDP) hook points in the Linux kernel network stack. The red loop indicates the 'shortcut' that is utilized by the fast-relay. TC hook allows redirection directly to egress while XDP hook is only available for ingress processing.
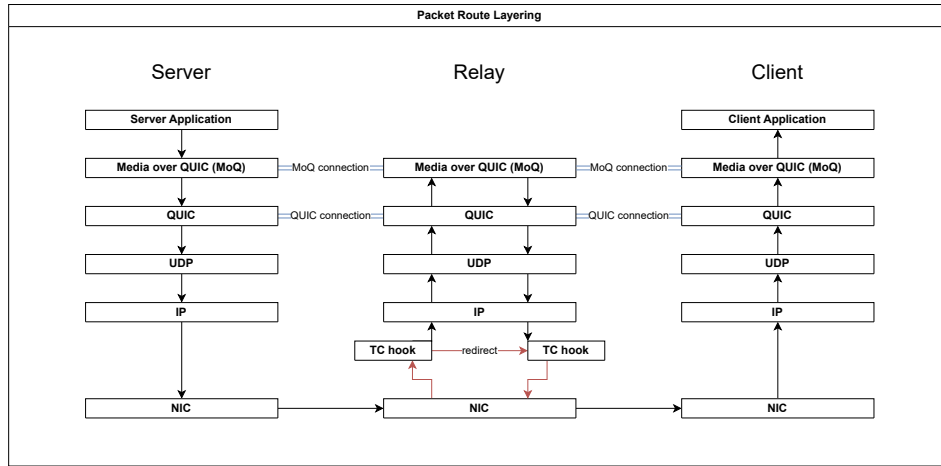
2.2.5 eBPF and Fast-Relays

TODO

Figure 2.2: Conventional layers of a network stack for client, server and relay. The red loop indicates again the 'short-cut' that is utilized by the fast-relay and based on eBPF packet-forwarding. This avoids the need for the packet to traverse the entire network stack of the relay up to the userspace.

2.3 Media over QUIC (MoQ)

On the application layer we will use the Media over QUIC (MoQ) protocol which is as of summer 2024 still being in the process of standardization by the IETF. MoQ targets live-streaming and real-time collaboration applications like Zoom, Microsoft Teams, or Google Meet. It is built on top of the QUIC protocol with the possibility of using WebTransport for browser support. A general publisher/subscriber model is used and the draft tries to combine performant approaches from protocols like RTP (for real-time features) and HLS/DASH (for scalability).

Figure 2.3: Rought MoQ relay architecture.

2.3.1 Solving Scaling versus Latency

For a long time now there have been two different camps with regards to media-data-transmission-protocols and -setups. One is heavily focused on low latency while the other is aiming for high scalability. Systems of the former kind include real-time collaboration tools like aforementioned Zoom, Teams, or Meet. The latter ones are often huge platforms like Twitch, YouTube or Netflix which need to reach millions of users at the same time. The one thing both have in common is that it turns out to be difficult to incorporate both low latency and high scalability into the system at the same time. The MoQ protocol tries to solve this by providing a setup that is both low-latency and highly scalable. To achieve this it supports performance enhancing approaches like relay caching or support for adaptive rate mechanisms.

2.3.2 Design of a MoQ Relay

The charter for the IETF working group describes what MoQ, and therefore also a relay that wants to meet the MoQ requirements, needs to support. These requirements for the publication- and distribution-setup mention the support of multiple formats, dynamic rate adaption mechanisms (e.g. used for congestion handling) as well as cache-friendly mechanisms.

Figure 2.3 gives a visualization of the rough architecture of a MoQ relay. It hints at key components like the relay level cache and the congestion handling mechanism. What one can also infer is the place of MoQ in the OSI-model namely at the application layer which itself builds on top of lower level protocols like QUIC, UDP, IP and Ethernet. In figure 2.4 and figure 2.5 one can see a comparison between the delayed fan-out of the media content caused by the MoQ relay and a more traditional setup. The former is able to reduce the traffic through a network by making multiple transmissions of the same data between server and relay obsolete.

This caching mechanism fits quite naturally into our proposed eBPF setup since we
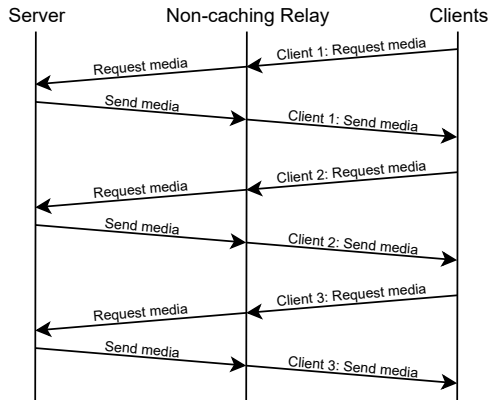
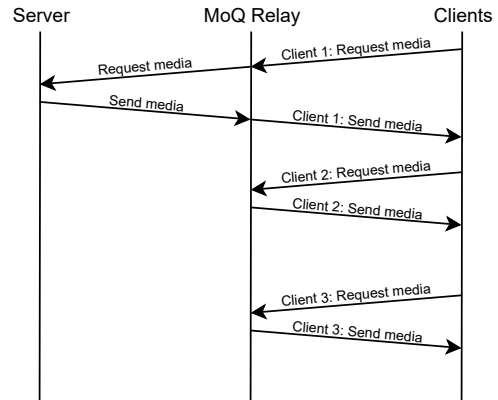Figure 2.4: Multiple data transmissions for different clients.

Figure 2.5: Only single data transmission towards the relay.

will need to communicate packet data between kernel- and user-space anyway to keep the QUIC library in a consistent state. In addition to that the congestion handling functionality of the MoQ relay can also be integrated fairly easy within eBPF. This is because packet dropping is ultimately one of the main use cases of plain eBPF programs and as such easy to implement. In a later section will go into detail on how the eBPF setup actually uses mechanisms like priority fields for this and how those meet the standard specifications.

2.4 Adaptive Bitrate Streaming

Multimedia streaming is a big part of the internet and many optimizations have been developed to improve the quality of service for the end-users. This includes considering (in real-time) parts of the clients connection state, such as available bandwith, and adapting the rate at which a server sends data. Such a process is called 'Adaptive Bitrate Streaming' and is employed in many of todays streaming setups. An example setup can be seen in Figure 2.6 where within the content delivery network multiple streams with different resolutions exist and the edge server that manages the connection to the client can switch between those streams based on the clients connection state. Youtube and Netflix are examples where, although more complex, similar setups are used to provide a better user experience.

The way the example implementation of fast-realys in this thesis is set up, is that the video server will encode into each packet its priority. For example I-frames have a high priority while P-frames have a lower priority. The relay then can decide not to forward
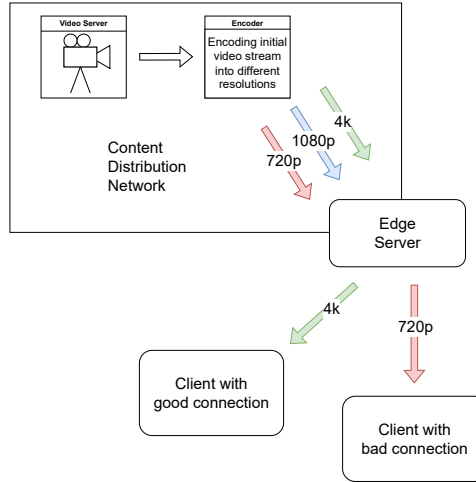
Figure 2.6: A streaming server might send multiple streams with different resolutions to allow adapting to a users bitrate.

certain packets to a client if the client is experiencing congestion.

TODO more specifics on how client state is found out.

2.5 Related Work

2.5.1 eQUIC Gateway

There have been previous publications on making QUIC more efficient by using BPF programs such as [PVV21], where a BPF program is used together with the Linux eXpress Data Path (XDP) to filter packets based on information provided by the userspace. This approach provided significant performance improvements with an increase of throughput by almost a third and a reduction of CPU time consumption caused by filtering packets by more than 25%. This shows that a setup leveraging Linux kernel features such as BPF has a lot of potential to improve current infrastructure.

2.5.2 Kernel Bypass

Another interesting approach which follows a similar idea of speeding up packet processing by avoiding the Linux network stack is [Tyu22]. The difference in this work is that DPDK is used to bypass the network stack to then process packets in userspace instead of using BPF programs like this thesis does. This, for example, offers more

flexibility as the userspace program is not as limited (e.g. by the BPF verifier) as the BPF program but might also lead to slightly more system calls, especially in the setup of a system, when user- and kernel-space need to communicate.

2.5.3 Priority drop

The idea of dropping packets based on their priority to adapt a connection in a congestion event has also been around for a while. [KWF03] explores this in more detail. Mainly improvements like a more tailorable congestion handling than the sole usage of discrete video quality levels as well as an improvement to, potentially randomized, frame dropping are discussed. This thesis, similar to [KWF03], will not focus on how the priority for packets is determined but rather on how those marked packets are handled. For this it is assumed that a higher level protocol has corectly determined the packet priorities and can handle the drop of packets with lower priority in case of limited bandwidth.

3 Fast-Relays

3.1 QUIC Adaptions

As was already mentioned in the previous chapter, our setup requires some adaptations
to the quic-go library. One initial change that was necessary was to turn off packet en-
and decryption, happening within quic-go. Given that we operate on the QUIC-header
data within the eBPF-program we need access to fields that are encypted using QUICs
header-protection. For obvoius reasons sending unencrypted packets is not something
that would be wanted in a production environment but for our setup it is required
since no fitting hardware offload was available at time of writing that would have
allowed us to 'push down' en- and decryption onto a smartNIC. Given that such a
hardware offload is added in the future, the en- and decryption can be turned on again
which makes this change more of a temporary solution to show the feasibility of our
approach.
Another type of change that we needed to introduce into the quic-go library is caused
by connection state management. We essentially added support for communication
with the eBPF-program by using an approach similar to C-style function pointers.
On multiple locations we added conditional function calls like the one depicted in
Listing 3.1. The function that is called here will be defined by the developer of the relay
and therefore allow for customizability without the need for changing the library itself.

```
/* Function pointer call within actual quic-go code */
if packet_setting.ConnectionUpdateBPFHandler != nil /* && potentially other
    conditions */ {
    packet_setting.ConnectionUpdateBPFHandler(connID.Bytes(),
        uint8(connID.Len()), p.connection)
}
```

Listing 3.1: An example of a function-pointer addition to the quic-go library.

```
/* Function pointer signature definition within additional config file */
ConnectionUpdateBPFHandler func(id []byte, l uint8, conn QuicConnection) = nil
```

Listing 3.2: Only the signature will be defined within the library itself.

The definition of the function that the developer of the relay wished to be executed at the specifically defined points will be defined locally in the relay code and provided to the configuration of the quic-go library. An example of how this could look like is shown in Listing 3.3.

```
/* Definition of the function within the local relay code */
func UpdateConnectionId(id []byte, l uint8, conn
    packet_setting.QuicConnection) {
    /* handle the connection update by interacting with the eBPF-program */
}

/* Providing the function to the quic-go library */
func main() {
    /* ... */
    packet_setting.ConnectionUpdateBPFHandler = common.UpdateConnectionId
    /* ... */
}
```

Listing 3.3: An example of how the addition looks on the relay side.

The need for these additions arises since the eBPF-program works with its own copy of the current state of a connection. This, for example, includes the connection ID that will be used when changing the packet header before sending it out. Since a connection ID can change, i.e. be updated or retired, during the lifetime of a connection we need a way to inform the eBPF-program to no longer use outdated state-information. These function-pointer style additions provide a minimal way of adding such functionality without limiting flexibility or adding too much application specific code to the library itself as it would be the case if the library would access the eBPF-Maps directly.

3.2 eBPF Setup

3.2.1 Different BPF Programs

In order to allow the relay to forward packets independently of the userspace, we need to equip the relay with three BPF programs as seem in figure. Those three programs are

- a program that handles incoming traffic **from** the clients (client ingress),

- a program that handles outgoing traffic **to** the clients (client egress) and

- a program that handles incoming traffic **from** the video server (server ingress).

Figure 3.1: The relay has to be equipped with three BPF programs.

Their responsibilities then are

- handling the initial registration of new clients and storing their information such as MAC addresses in a BPF map,

- intercepting the packets from the video server, duplicating and redirecting them to the egress program (as well as sending one unaltered packet to userspace for state management purposes),

- receiving the redirected packets at egress, altering them using the client specific data, deciding (based on packet priority and client congestion) if a packets should be dropped or sent, storing info on sent out packets for future congestion control purposes and finally sending them out to the clients.

This setup allows us to separate any state management and congestion control from the actual packet forwarding and thus makes leaving out any immediate userspace processing possible.

Following is a more detailed description of the responsibilities of each of the three programs.

Client Ingress

TODO

Server Ingress

TODO

Client Egress

The client egress program sees every packet that leaves the relay. This includes packets that have been redirected by the ingress program as well as packets that have been generated by the relay itself. Since the QUIC protocol works with packet numbers for a given connection it is necessary for the egress program to make sure the forwarded packets together with the userspace packets provide a consistent state. For this the egress program maintains its own packet number counter for each connection. That way only one counter has to be maintained and race conditions can be avoided. However, this also means that the packets sent by the userspace are likely to have a different packet number than the one chosen by the QUIC library. This might lead to inconsistencies again but can be avoided by not storing a packet from userspace right away in the packet history but only once the BPF has stored it, along with the changed packet number, in the map used for packet registration. This initially gives a brief window where a packet was sent out but is not saved in the history of the QUIC library but once the packet is then processed by the userspace routine handling the registration, any incoming ACKs for this packet can be processed correctly. TODO

3.2.2 Packet Registration

In order to make the congestion control algorithm that is running in userspace usable we need to inform the QUIC library about the forwarded packets. This again happens via BPF maps and a separate go routine that continuously polls new entries in the map and processes them. Entries are then added to the packet history to allow the receipt of ACKs. Besides that, the congestion control algorithm will be informed about the forwarded packet in order to be able to react to potential congestion events.

TODO: mention

3.3 User Space Avoidance

TODO

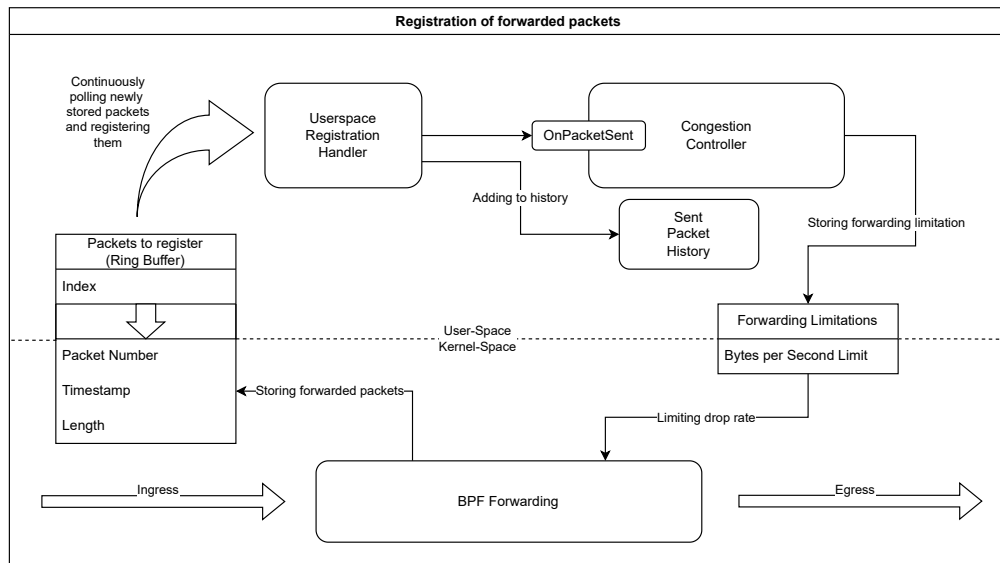## 3.4 Packet Filtering and Dropping

TODO

Figure 3.2: Internal setup for registering forwarded packets as well as incorporating forwarding limitations for the BPF program.

3.5 Client Congestion

TODO

3.6 Subscription and State Management

TODO

3.7 Relay Caching

TODO

3.8 Compatibility

TODO

Figure 3.3: TODO (also mention that the check for previous translation is needed since >1 packets per unistream are possible)

3.9 Source Code Repositories

For the development of the relay and the eBPF programs, we have come up with the following repositories:

- **Fast-Relay** [Pfe24b]: This is the main repository providing the eBPF program implementations as well as examples of server, relay and client implementations in Go.

- **Quic-Go Adaptation** [Pfe24a]: This repository is a fork of the QUIC library "quic-go" [See24] and provides a plain Go implementation of the QUIC protocol. For our thesis we needed to make some adaptations to the library to support some hook points for separate functions which should be specifically designed to handle the underlying eBPF setup with its eBPF-Map usage.

- **MoQ-Transport Adaptation** [PE24]: This repository is a fork of the "MoQ-Transport" [Int24a] protocol repository and provides some needed adaptations to our examples. One such adaptation is that the server needs to support a categorization of payloads into different priorities in order for the eBPF program to be able to deliberalty drop packets in case of congestion. Getting these priorities could be as simple as differentiating only between I- and P-frames in a video stream or more complex based on the needs of the application and the wanted granularity of the congestion control.

4 Testing

4.1 Setups

TODO

### 4.1.1 Namespace Environment for Local Testing and Development

TODO

4.1.2 Physical Server Setup for Real-World Testing

TODO

```
Logical setup:
                                interface: veth1                     interface: veth3
                                ip: 192.168.1.2/24                   ip: 192.168.1.4/24
              interface: veth0          \           interface: veth2          \
              ip: 192.168.1.1/24         \          ip: 192.168.1.3/24         \
 _____    /                          \  _____   /                        \  _____
|          |  /                            \|       | /                          \|          |
|  Server  |-------> veth0@veth1 ------->|  Relay |-------> veth2@veth3 ------->|  Client  |
|_____|<------- veth1@veth0 <-------|_____|<------- veth3@veth2 <-------|_____|
```

```
Actual setup (using bridge and namespaces):
 ----------------------      ---------------------------------------      ----------------------
|   Server namespace   |    |              Relay namespace          |    |   Client namespace   |
|    -----------       |    |    -------------    -------------     |    |    -----------       |
|   |192.168.10.1|     |    |   | 192.168.10.2 | | 192.168.11.2 |   |    |   |192.168.11.1|     |
|____|___veth0____|____|    |___|____veth1_____|_|____veth2_____|___|    |____|____veth3___|____|
          \                           /              \                            /
           \                         /                \                          /
            \                       /                  \                        /
             \                     /                    \                      /
              \ _____ /_                    __\ _____ /
              /veth0-br|    |veth1-br\                 /veth2-br|    |veth3-br\
              |           ___|                         |___               |
              \_____v-net-0_____|NAT/                \NAT|_____v-net-1____/
                    /              \                    /               \
              ip: 192.168.10.10     \                  /      ip: 192.168.11.10
              net: 192.168.10.0/24   \  _____   /       net: 192.168.11.0/24
                                      / _____  \
                                     |    enp1s0f0    |
                                      \ _____  /
                                           |
                                           |
                                      (INTERNET)
```

4.2 Testing

TODO

4.3 Results

TODO

4.3.1 Delay Reduction of BPF Forwarding

When looking at the impact of eBPF-Forwarding on the delays of packets, we can see that in the namespace environment the delay of a single packet is decreased by around 100 µs when compared to the userspace forwarding. This is shown in Figure 4.1 and consideres the simplest case of the userspace program where it only forwards directly to one connection without the need for much additional computation. Given that the userspace can have arbitrary complex connection management this delay improvement can likely be even bigger in a more sophisticated userspace setup.
Another thing the figure shows is that the delay has a smaller variance due to the fact that the eBPF program path is somewhat similar for each packet whereas, in contrast,

the userspace path can have buffers, queues, or similar that lead to a higher difference in processing time between packets. This effect however might be less observable in a real world scenario due to the ubiquitous network jitter which was influencial in the used namespace environment.


Figure 4.1: By avoiding userspace when processing the 1-RTT packets that contain the payload the delay of a single packet can be reduced. The longer-delay-operations are either handled directly in the eBPF program (e.g. the case for deciding where to redirect a packet to) or handled after the packet has already been sent out (e.g. the case for registering a packet such that the QUIC library knows about it).

Figure 4.2: Server namespace CPU usage comparison.

Figure 4.3: Relay namespace CPU usage comparison.

Figure 4.4: Client namespace CPU usage comparison.

Table 4.1: Table for CPU usage of relay Go processes

flat	flat%	sum%	cum	cum%	cause
440ms	28.39%	28.39%	440ms	28.39%	runtime/internal/syscall.Syscall6
140ms	9.03%	37.42%	140ms	9.03%	runtime.futex
120ms	7.74%	45.16%	120ms	7.74%	runtime.cgocall
40ms	2.58%	47.74%	40ms	2.58%	runtime.write1
30ms	1.94%	49.68%	60ms	3.87%	runtime.checkTimers
30ms	1.94%	51.61%	50ms	3.23%	runtime.mapaccess2
20ms	1.29%	52.90%	70ms	4.52%	.../cilium/ebpf/internal/unix.Syscall
20ms	1.29%	54.19%	20ms	1.29%	net.IP.String
20ms	1.29%	55.48%	20ms	1.29%	runtime.casgstatus
20ms	1.29%	56.77%	20ms	1.29%	runtime.duffcopy

5 Future Work

5.1 Hardware Offload

This thesis heavily relies on the fact that the relay can access certain fields (e.g. packet numbers) of the packet, which are generally not accessible without prior decryption. In the current setup, this is made possible by turning off encryption alltogether but to be of any use in a real-world scenario, the encryption of incoming and the decryption of outgoing packets would need to be pushed down below the lowest used BPF hook point in the stack. This means that a hardware offload of encryption and decryption similar to what is done for TCP/IP checksums would be necessary.

Once compatible SmartNIC offload implementations are available one can, besides en- and decryption, also offload the BPF program itself. This then would provide another way of accelerating performance.

Some previous work in this direction has already been done since at least 2019 [Yan+20] but for the purpose of this thesis we did not find any suitable open-source implementation that would allow incorporation into our fast-relay example implementation.

5.2 Compatibility Expansion

This thesis used the QUIC protocol together with media over QUIC (MoQ) to demonstrate how fast-relays that circumvent userspace by utilizing BPF programs can be designed. However, generally speaking the design of fast-relays is not limited to any of these protocols and could be expanded given modifications to necessary fields are possible (i.e. not prevented by encryption) and there is a way to encode the priority of a packet in the packet itself. The latter point could always be realized by using part of the payload which forces a deeper packet inspection within the BPF program but avoids the need to fit the priority into the header of an existing protocol.

6 Conclusion

TODO

List of Figures

List of Tables

Bibliography

[DB19]   M. Deval and G. Bowers. "Technologies for Accelerated QUIC Packet Processing with Hardware Offloads." Patent EP3541044A1. European Patent application made by Intel Corporation. Sept. 2019.

[Fou]   L. Foundation. *eBPF Verifier*. Accessed: July 13th, 2024. URL: https://docs.kernel.org/bpf/verifier.html.

[Int24a]   Internet-Engineering-Task-Force. *Media over QUIC Transport*. Accessed: May 26th, 2024. 2024. URL: https://datatracker.ietf.org/doc/draft-ietf-moq-transport/.

[Int24b]   Internet-Engineering-Task-Force. *QUIC: A UDP-Based Multiplexed and Secure Transport*. Accessed: May 26th, 2024. 2024. URL: https://datatracker.ietf.org/doc/html/rfc9000.

[JC20]   M. Joras and Y. Chi. *How Facebook is bringing QUIC to billions*. Accessed: May 26th, 2024. Oct. 2020. URL: https://engineering.fb.com/2020/10/21/networking-traffic/how-facebook-is-bringing-quic-to-billions/.

[KWF03]   C. Krasic, J. Walpole, and W.-c. Feng. *Quality-Adaptive Media Streaming by Priority Drop*. June 2003. DOI: 10.1145/776322.776341.

[Lan+17]   A. Langley, A. Riddoch, A. Wilk, A. Vicente, C. Krasic, D. Zhang, F. Yang, F. Kouranov, I. Swett, J. Iyengar, J. Bailey, J. Dorfman, J. Roskind, J. Kulik, P. Westin, R. Tenneti, R. Shade, R. Hamilton, V. Vasiliev, W.-T. Chang, and Z. Shi. "The QUIC Transport Protocol: Design and Internet-Scale Deployment." In: *SIGCOMM '17: Proceedings of the Conference of the ACM Special Interest Group on Data Communication* (Aug. 2017), pp. 183–196. DOI: 10.1145/3098822.3098842.

[Lie18]   E. Liebetrau. *How Google's QUIC Protocol Impacts Network Security and Reporting*. Accessed: May 26th, 2024. June 2018. URL: https://www.fastvue.co/fastvue/blog/googles-quic-protocols-security-and-reporting-implications/.

[PE24]   D. Pfeifer and M. Engelbart. *Media over QUIC Transport*. Accessed: May 26th, 2024. 2024. URL: https://github.com/danielpfeifer02/priority-moqtransport.git.

[Pfe24a]   D. Pfeifer. *A QUIC implementation in pure Go*. Accessed: May 26th, 2024. 2024. URL: https://github.com/danielpfeifer02/quic-go-prio-packs.git.

[Pfe24b]   D. Pfeifer. *Adaptive Media over QUIC Transport*. Accessed: May 26th, 2024. 2024. URL: https://github.com/danielpfeifer02/Adaptive_MoQ.git.

[Pfe24c]   D. Pfeifer. *eBPF-Assisted Relays for Multimedia Streaming*. Accessed: May 26th, 2024. 2024. URL: https://github.com/danielpfeifer02/tum-thesis-bsc-fast-relays.git.

[PVV21]   G. Pantuza, M. A. M. Vieira, and L. F. M. Vieira. *eQUIC Gateway: Maximizing QUIC Throughput using a Gateway Service based on eBPF + XDP*. 2021. DOI: 10.1109/ISCC53001.2021.9631262.

[Rod24]   L. Rodriguez. *The Future of the Internet is here: QUIC Protocol and HTTP/3*. Accessed: May 26th, 2024. Jan. 2024. URL: https://medium.com/@luisrodri/the-future-of-the-internet-is-here-quic-protocol-and-http-3-d7061adf424f.

[See24]   M. Seemann. *A QUIC implementation in pure Go*. Accessed: May 26th, 2024. 2024. URL: https://github.com/quic-go/quic-go.git.

[Sta94]   I. O. for Standardization. *ISO/IEC 7498-1:1994*. Accessed: July 3rd, 2024. 1994. URL: https://www.iso.org/obp/ui/en/#iso:std:iso-iec:7498:-1:ed-1:v2:en.

[Tyu22]   N. Tyunyayev. *Improving the performance of picoquic by bypassing the Linux Kernel with DPDK*. Master's Thesis. 2022.

[W3T24]   W3Techs.com. *Usage statistics of QUIC for websites*. Accessed: May 26th, 2024. May 2024. URL: https://w3techs.com/technologies/details/ce-quic.

[Yan+20]   X. Yang, L. Eggert, J. Ott, S. Uhlig, Z. Sun, and G. Antichi. "Making QUIC Quicker With NIC Offload." In: *EPIQ '20: Proceedings of the Workshop on the Evolution, Performance, and Interoperability of QUIC* (Aug. 2020), pp. 21–27. DOI: 10.1145/3405796.3405827.