



TUM SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

eBPF-Assisted Relays for Multimedia Streaming

Daniel Alexander Antonius Pfeifer



TUM SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

eBPF-Assisted Relays for Multimedia Streaming

eBPF-Unterstützung für Multimedia-Streaming-Netznoten

Author:	Daniel Alexander Antonius Pfeifer
Supervisor:	Prof. Dr.-Ing. Jörg Ott
Advisor:	Mathis Engelbart, M.Sc.
Submission Date:	15.08.2024

I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

Munich, 15.08.2024

Daniel Alexander Antonius Pfeifer

Acknowledgments

Abstract

In this thesis we propose a new relay setup for multimedia streaming that allows for avoidance of userspace processing by utilizing BPF programs in the Linux kernel. In a sample implementation we demonstrate the feasibility of this approach by designing a relay that is capable of forwarding packets between a server-side and a client-side QUIC connection while still being able to do adaptive bitrate streaming based on client congestion.

We show that this approach saves processing time and reduces latency compared to userspace processing with the relay still adhering to specifications of the QUIC standard and the 'Media over QUIC' (MoQ) draft. One limitation that is not addressed in depth in this thesis is the need for a de- and encryption hardware offload onto a SmartNIC to allow the BPF program to access the packet payload without any restrictions. Since the QUIC standard is still fairly new we are confident that a solution for a potential hardware offload will be found in future research.

Contents

Acknowledgments	iii
Abstract	iv
1 Introduction	1
1.1 Research Question	1
1.2 Scope	2
1.3 Structure of this Thesis	3
2 Background and Related Work	4
2.1 QUIC	4
2.1.1 Connections and Streams	4
2.1.2 quic-go	5
2.1.3 QUIC's Importance to Fast-Relays	5
2.2 eBPF	5
2.2.1 eBPF Hook Points	6
2.2.2 eBPF Verifier	7
2.2.3 Important eBPF Concepts	7
2.3 Media over QUIC (MoQ)	8
2.3.1 Solving Scaling versus Latency	8
2.3.2 Design of a MoQ Relay	9
2.3.3 moqtransport	10
2.4 Adaptive Bitrate Streaming	10
2.4.1 Mechanisms and Ideas	11
2.4.2 Implications on Streaming Setup	11
2.5 Related Work	11
2.5.1 eQUIC Gateway	11
2.5.2 Kernel Bypass	11
2.5.3 Priority drop	12
3 Fast-Relays	13
3.1 QUIC Adaptions	13
3.1.1 Function-Pointer Style Additions	14
3.1.2 Direct Changes to the Library	15

3.2	eBPF Setup	17
3.2.1	Different eBPF Programs	17
3.2.2	Packet Registration	23
3.2.3	Retransmissions of Forwarded Packets	24
3.3	Userspace Synchronization	26
3.3.1	Subscription and State Management	26
3.3.2	Relay Caching	26
3.4	Congestion Considerations	27
3.4.1	Client Congestion	27
3.4.2	Packet Filtering and Dropping	27
3.5	Integration and Prototype	28
3.5.1	Compatibility	29
3.5.2	Source Code Repositories	29
4	Testing	30
4.1	Setups	30
4.1.1	Local Environment for Testing and Development	30
4.1.2	Physical Server Setup for Real-World Testing	31
4.2	Testing and Results	31
4.2.1	Delay Reduction of eBPF Forwarding	31
4.2.2	CPU Utilization Comparison	33
5	Future Work	35
5.1	Hardware Offload	35
5.2	Compatibility Expansion	35
6	Conclusion	36
	List of Figures	37
	List of Tables	38
	Bibliography	39

1 Introduction

The fact that online streaming tends to be slower than cable TV is likely something most people have already experienced first-hand. Live sports events, music shows, or news broadcasts arrive an order of seconds later when streamed compared to using traditional cable connections. Despite this delay not necessarily being a deal-breaker for most people, designing networks that tighten the gap between cable and streaming is still worthwhile. With very optimized and fast networks already in place, we are at a point where providing such faster information delivery is highly non-trivial. To do that, we have even gone as far as developing entirely new standards, such as QUIC. Those new standards aim to improve the shortcomings of some of the most fundamental protocols of the Internet, one of them being TCP, which have been around more than 40 years.

Besides introducing new protocols, one could also look at existing setups and figure out how to trade some generality for a smaller delay when handling data. The ISO/OSI model, which is a foundational concept in networking, provides “a common basis for the coordination of standards development for the purpose of systems interconnection” [Sta94]. As one can imagine, such a “common basis”, even though convenient for large-scale systems, can cause unnecessary overhead. In some cases, additional speed-ups can be achieved by using more application-specific approaches. This thesis will consider one such case and explore the possibilities of avoiding or delaying specific processing steps of the ISO/OSI model to increase the overall speed of a network relay.

1.1 Research Question

As already mentioned above, the usage of application-specific approaches in networking allows for a reduction in latency. In this thesis, we will consider a media streaming scenario that runs on top of QUIC by using the “Media over QUIC” (MoQ) transport protocol [Int24a]. The central question we will try to answer in this thesis will then be:

How can we improve the performance of a relay in a media streaming scenario by using eBPF technology?

By using eBPF technology together with kernel hook points provided by the Linux-Kernel, we will try to find a setup that improves relay performance using eBPF programs

that handle basic relay capabilities, such as packet forwarding and congestion control. Since the QUIC protocol is designed to handle a significant portion of its workload in userspace, we look into possibilities of delaying any userspace processing until **after** the packet has been forwarded to the client. This way, the raw delay that the packet experiences from the initial media server to the client could be reduced. However, since QUIC is a connection-oriented protocol, we need to make sure that the QUIC connection state stays coherent despite the additional processing steps done by the eBPF program. We will investigate which additional processing steps are needed in our case, how they compare to challenges when expanding our approach to other protocols, and how they can be implemented in an eBPF program. Therefore, more specific sub-questions we try to answer are:

1. *How can we avoid the need to direct a packet through userspace?*
2. *How to handle the fact that packets are heavily encrypted?*
3. *What communication between userspace and eBPF program is necessary to keep coherency?*
4. *How can our approach be generalized to other protocols?*

1.2 Scope

The main improvement this thesis aims to achieve is shortening a packet's critical path from a media server to a client. This will be done by avoiding the immediate need for a packet traversal up the network stack to the application layer. Instead, any communication with the application layer will happen in a delayed fashion (after the packet is sent) by utilizing eBPF-Maps for storing any necessary (meta-) information. The main reason this communication between userspace and eBPF-program is required is that relays in MoQ are an application layer concept. That means the QUIC connections from relay to server and from relay to client will be different, and the packets that have been eBPF-forwarded to egress directly will need changes in their header data to match the state of the outgoing client connection.

This approach is highly dependent on the used standards and protocols. This thesis will operate on top of the QUIC protocol [Int24b] and the "Media over QUIC" (MoQ) transport protocol [Int24a]. For the application layer, the quic-go library [See24] will provide the implementation and any additional (non-eBPF) program will also be written in Go. Since the setup is dependent on retrieving data from eBPF-Maps, the QUIC library providing the implementation will need some adaptations. We will mainly introduce simple function pointer style additions that allow the adapted library to be

run both with and without the eBPF setup. The developer of the relay will then also have more freedom to set up the eBPF part of the relay as they see fit since the Go code that will interact with eBPF parts will also have to be provided by said developer.

Additionally, we will run a performance analysis on our implementation of the relay to confirm the potential this approach has. These performance tests will look at the raw delay speedup as well as the impact on CPU utilization this setup has. All the tests will be done in a lab-like environment to isolate the performance changes as best as possible from any outside noise. Most payloads used will only contain dummy data since our approach does not interfere with payload contents and there is no need for creating and using actual media stream data.

Despite our approach only considering QUIC and MoQ, we will argue that the general idea of our setup will be independent of any of these protocols and can be changed to fit one's needs.

With this, we will provide answers to the research questions regarding packet-redirection, communication between userspace and eBPF as well as setup-generalization. Regarding the question on how to handle the encryption of the packets, we will not focus on this since we did not find a suitable hardware offload that would have allowed for en- and decryption after and before the used eBPF hook points, respectively. Instead, we will emulate this behavior by turning off the encryption in the QUIC library itself, which will provide a similar result.

1.3 Structure of this Thesis

In chapter 2, this thesis will provide some overview of used technologies and related ideas. Section 2.1 will give an introduction to the QUIC protocol and its main features and section 2.2 will provide an overview of eBPF technology together with features related to our approach. Section 2.3 will introduce the 'Media over QUIC' (MoQ) transport protocol which will be used for our application-level relay setup. After that section 2.4 will explain the ideas and challenges of adaptive bitrate streaming while section 2.5 will mention some work related to the aforementioned topics. What will follow in chapter 3 is a detailed description of the setup that allowed us to improve relay performance. We will look at the adaptations to the used QUIC library in section 3.1 as well as our eBPF setup in 3.2. Besides those two, we will also look at some more specific details and challenges in the subsequent sections. In chapter 4, we will then provide a basic performance analysis of our setup to show current improvements and limitations. Finally, we will conclude with a summary together with some ideas for future work in this field in chapters 5 and 6.

2 Background and Related Work

2.1 QUIC

Many fundamental internet protocols still used today have been around for a very long time. For example, the Transmission Control Protocol (TCP) has been used as the backbone of the internet for more than 40 years. It has been designed to be reliable and to provide a connection-oriented way of transmitting data, but the modern environment of the internet with needs like lower latency, better multiplexing, or improved security makes it hard for TCP to keep up. Limitations in the design and resulting issues like head-of-line blocking have raised demand for a newly designed protocol that can keep up with the modern internet. All of these issues, paired with the want for a more flexible development cycle led to new creations. QUIC, which started off as the ‘Quick UDP Internet Connections’ protocol and has since been standardized by the IETF, with QUIC now being its own trademark, is a transport layer protocol built on top of UDP that is designed to be reliable, cryptographically secure and more performant than TCP. QUIC, partly because it operates both in user- and kernel-space, has been designed to allow for a more rapid deployment cycle than TCP. Similar to TCP, it is a connection-based protocol that uses TLS for encryption [Lan+17]. Back in 2018, QUIC was already the default protocol for the Google Chrome browser, which, at the time, made up 60% of the web browser market [Lie18]. A little over two years later, Facebook, now Meta, was using QUIC for more than 75% of their internet traffic which led to improvements regarding request errors, tail latency, and header size [JC20]. As of May 2024, QUIC already made up 8.0% of all internet traffic with support from pretty much every major browser [W3T24; Rod24]. With big players like Google, Meta, Microsoft, or YouTube putting emphasis on using QUIC to improve their services, this number will likely increase even further.

2.1.1 Connections and Streams

Since QUIC is a connection-based protocol, some initial overhead to establish a connection is needed. However, the design incorporates features that aim for an efficient way of establishing connections, e.g. by using 0-RTT (zero round-trip-time) handshakes. Latency improvements like the 0-RTT handshake, however, come at the cost of security since that opens the door for replay attacks. Another part where QUIC tries to optimize connection management is the use of streams. Streams are designed to be lightweight and can be opened without the need for a handshake. This goes as far as one single

packet being able to open a new stream, transferring stream data as well as the closing the stream again. This allows for new techniques to improve data transmission and will also be part of the fast-relay setup in this thesis. Aside from streams, apparent since QUIC is based on UDP, it is also possible to send data via unreliable datagrams. This further improves the versatility of the protocol and allows for new ways of optimizing data transmission.

2.1.2 quic-go

There are many implementations of the QUIC protocol available, providing libraries for a lot of today's most popular programming languages. The implementation we settled on for this thesis is the quic-go library which provides a pure Go approach to implementing the standards of RFC-9000, RFC-9221 as well as some others which are not important for our usecase. However, since we need some special behavior of the userspace part of QUIC, we will introduce some modifications into quic-go. Those modifications will be explained further in section 3.1.

2.1.3 QUIC's Importance to Fast-Relays

The QUIC protocol will be a fundamental part of the fast-relay setup in this thesis, yet the ideas used to make relays faster is not limited to QUIC and can be extended to other protocols as well. QUIC is chosen as an example protocol due to its increasing popularity, which offers big potential in early adoption and deployment of fast-relays. Also, the easy incorporation of changes into libraries providing RFC implementations makes it a good starting point for experimenting with what can and cannot be done regarding our research questions. This includes the possibility of neglecting the difficulties that the heavy encryption of QUIC brings with it by just turning off the related functionality.

2.2 eBPF

In 1992 a technology called 'Berkeley Packet Filter' (BPF) was introduced into the Unix kernel. By using BPF, it is possible to attach a small BPF-program to some pre-defined hook points in the network stack of the kernel and filter packets there in a stateless manner. This provided more efficiency since the packets did not need to be copied into userspace anymore but could be processed directly in the kernel. A need for better tracing capabilities of the Linux kernel then led to the development of an extended version of BPF called "eBPF" which was introduced in 2014 and heavily influenced by a tracing tool called "dtrace" [Tig].

2.2.1 eBPF Hook Points

The Linux kernel offers several hook points to which eBPF-programs can be attached. There are two prominent ones that we considered for our suggested setup. The first one allows one to access the Traffic Control (TC) subsystem, while the second one allows one to access the eXpress Data Path (XDP) subsystem.

XDP would generally provide a better performance since it is located lower on the network stack, namely directly in the NIC driver, than the TC-hook point, which is located in the link-layer. TC, on the other hand, offers a more versatile way of packet processing since the used `sk_buff` buffer object that contains the packet data provides access to metadata that is unavailable when using XDP and its `xdp_buff` buffer object. What ultimately led us to choose TC over XDP was, however, the fact that XDP only allows ingress packet processing, while TC allows both ingress and egress. That means that with XDP, we would not have been able to redirect packets to be handled at egress, which is crucial for the fast-relay setup we aim for.

Figure Figure 2.1 illustrates again the relative positions of the TC and XDP hook points in the network stack.

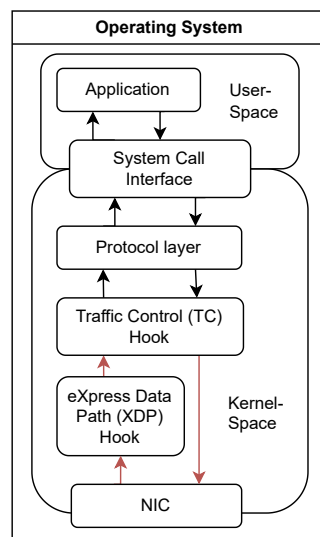


Figure 2.1: Abstracted view of Traffic Control (TC) and eXpress Data Path (XDP) hook points in the Linux kernel network stack. The red loop indicates the “short-cut” that is utilized by the fast-relay. TC hook allows redirection directly to egress, while XDP hook is only available for ingress processing.

2.2.2 eBPF Verifier

Since eBPF programs are executed in the kernel, it is quite obvious that extensive security checks need to be in place to ensure that the kernel does not experience problems like infinite loops, accesses to invalid memory locations, or other security related issues. This explains the existence of the so-called ‘BPF-verifier’ which inspects every BPF-program for its safety by simulating possible program paths, looking at the graph representation of the program and more [Fou]. This imposes some restrictions on the complexity of the programs that can be used within the kernel. In our case, this did not impose too many issues since we do not rely on very complex control structures.

2.2.3 Important eBPF Concepts

One of the most important concepts in eBPF, which we do use quite extensively, is the “eBPF-map”. Such a map boils down to a section in memory that is reserved for the eBPF-program and which can be used as a key-value store for arbitrary data. This part of memory can then also be accessed from userspace and thus provides the main way of communication between the eBPF-program and our application. When we define an eBPF-map, we can choose between different types and configure size, key-type, value-type, and the way the map is stored. An example of two eBPF-map definitions can be seen in Figure 2.2. It shows two different types of maps, a hash map, and a ring buffer, that are used in our fast-relay setup. Those and some other relevant map types are listed in table Table 2.1.

```
1 struct {  
2     __uint(type, BPF_MAP_TYPE_HASH);           // Hash map  
3     __type(key, struct client_info_key_t);     // Specific client key  
4     __type(value, uint32_t);                   // 32 bit id  
5     __uint(max_entries, MAX_CLIENTS);          // Maximum number of clients  
6     __uint(pinning, LIBBPF_PIN_BY_NAME);      // Pin by name to the tc filesystem  
7 } client_id SEC(".maps");  
8  
9 struct {  
10    __uint(type, BPF_MAP_TYPE_RINGBUF);        // Ring buffer  
11    __uint(max_entries, MAX_PACKET_EVENTS);    // Maximum number of packet events  
12    __uint(pinning, LIBBPF_PIN_BY_NAME);      // Pin by name to the tc filesystem  
13 } packet_events SEC(".maps");
```

Figure 2.2: Exemplary eBPF map definitions.

Type	Description
BPF_MAP_TYPE_HASH	A hash map where keys and values can be arbitrarily defined.
BPF_MAP_TYPE_PERCPU_HASH	A hash map with separate value slots for each CPU, providing improved performance in multi-core environments.
BPF_MAP_TYPE_ARRAY	An array map that allows random access to elements by index.
BPF_MAP_TYPE_PERCPU_ARRAY	An array map with separate value slots for each CPU, useful for per-CPU data storage.
BPF_MAP_TYPE_RINGBUF	A ring buffer for implementing high-performance data queues.

Table 2.1: Some eBPF map types. (defined in `/usr/include/linux/bpf.h`)

2.3 Media over QUIC (MoQ)

On the application layer, we will use the Media over QUIC (MoQ) protocol which is, as of summer 2024, still in the process of standardization by the IETF [Int23]. MoQ targets live-streaming and real-time collaboration applications like Zoom, Microsoft Teams, or Google Meet. It is built on top of the QUIC protocol with the possibility of using WebTransport for browser support. A general publisher/subscriber model is used, and the draft tries to combine performant approaches from protocols like RTP (for real-time features) and HLS/DASH (for scalability).

2.3.1 Solving Scaling versus Latency

For a long time now, there have been two different camps with regard to media-data-transmission-protocols and -setups. One is heavily focused on low latency, while the other aims for high scalability [Int24c]. Systems of the former kind include real-time collaboration tools like aforementioned Zoom, Teams, or Meet. The latter ones are often huge platforms like Twitch, YouTube, or Netflix, which need to reach millions of users at the same time. The one thing both have in common is that it turns out to be difficult to incorporate both low latency and high scalability into the system at the same time [Int24c]. The MoQ protocol tries to solve this by providing a setup that is both low-latency and highly scalable. To achieve this, it supports performance-enhancing approaches like relay caching or support for adaptive rate mechanisms.

2.3.2 Design of a MoQ Relay

The charter for the IETF working group [Int23] describes what MoQ, and therefore also a relay that wants to meet the MoQ requirements, needs to support. These requirements for the publication- and distribution-setup mention the support of multiple formats, dynamic rate adaption mechanisms (e.g. used for congestion handling) as well as cache-friendly mechanisms.




Figure 2.3: Rough MoQ relay architecture.

Figure 2.3 gives a visualization of the rough architecture of a MoQ relay. It hints at key components like the relay level cache and the congestion handling mechanism. What one can also infer is the place of MoQ in the OSI-model namely at the application layer, which itself builds on top of lower-level protocols like QUIC, UDP, IP, and Ethernet. In figure 2.4 and figure 2.5 one can see a comparison between the delayed fan-out of the media content caused by the MoQ relay and a more traditional setup. The former can reduce the traffic through a network by making multiple transmissions of the same data between server and relay obsolete.




Figure 2.4: Multiple data transmissions for different clients.

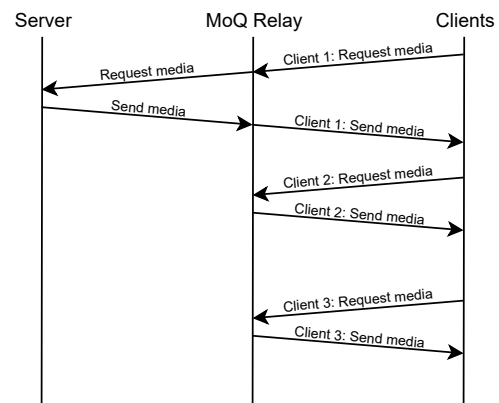


Figure 2.5: Only single data transmission towards the relay.

This caching mechanism fits quite naturally into our proposed eBPF setup since we will need to communicate packet data between kernel- and user-space anyway to keep the QUIC library in a consistent state. In addition to that, the congestion-handling functionality of the MoQ relay can also be integrated fairly easily within eBPF. This is because packet dropping is ultimately one of the main use cases of plain eBPF programs and, as such, easy to implement. In a later section we will go into detail on how the eBPF setup actually uses mechanisms like priority fields for this and how those meet the standard specifications.

2.3.3 moqtransport

To use the MoQ protocol in our setup, we will make use of a library that implements the MoQ transport protocol (moqtransport or MoQT) in Go [PE24]. The goal of MoQT is to define a media transport protocol that is operating on top of QUIC and WebTransport, which is providing the concrete designs that the charter of the MoQ working group is aiming for. This includes, for example, the actual structure of the different message types or error handling in case of a wrong state. As of writing this thesis, the MoQT draft is in its fifth version [Int24a].

2.4 Adaptive Bitrate Streaming

Multimedia streaming is a big part of the internet and many optimizations have been developed to improve the quality of service for the end-users. This includes considering (in real-time) parts of the clients' connection state, such as available bandwidth, and adapting the rate at which a server sends data. Such a process is called "Adaptive Bitrate Streaming" (ABS) and is employed in many of today's streaming setups.

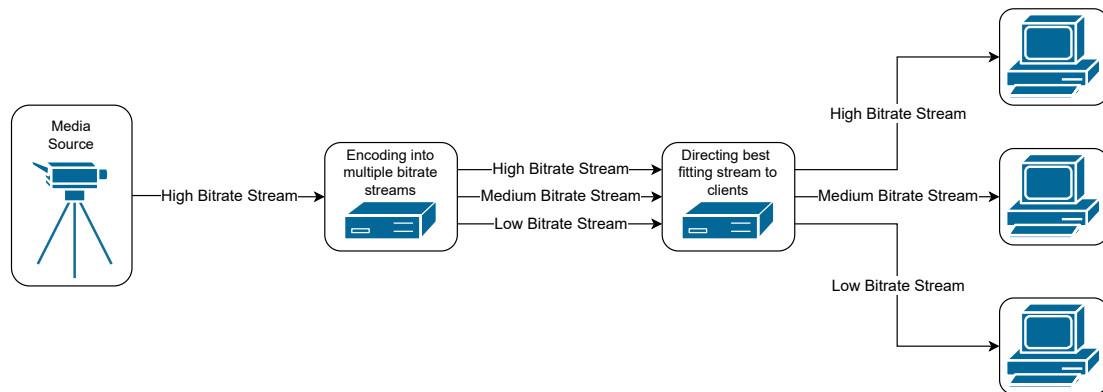


Figure 2.6: Behind the scenes multiple streams with different resolutions exist to allow adapting to a user's bitrate.

2.4.1 Mechanisms and Ideas

ABS implements the simple idea of changing the amount of data sent to a client based on the clients' connection quality. This means a client with a good connection gets a higher-quality stream than one with a bad connection. For this, the internal setup needs to be able to provide multiple streams with different resolutions and bitrates. An example setup can be seen in Figure 2.6 where an encoder creates streams with different resolutions and an edge server, that manages the connections to the clients, can switch between those streams based on the connection state of a specific client. YouTube and Netflix are examples where, although more complex, similar setups are used to provide a better user experience. The premise of such a setup is that in case of a bad connection, it is still preferable to be able to watch a stream continuously, even if that means cutting down on quality.

2.4.2 Implications on Streaming Setup

ABS forces a few changes to the way servers and relays are set up. First, as figure Figure 2.6 shows, there needs to be an "encoding" component that turns the single, high-quality stream that is received from the streaming source into multiple streams with different resolutions. Besides that, the relay needs to keep track of the clients' connection quality, e.g. using measurements like RTT, packet loss, etc., and decide which stream to forward to the client. Since the connection quality can change dynamically the relay also needs to continuously monitor each connection and potentially change the used stream. Besides encoding and monitoring aspects the system also has increased storage requirements since stream data will essentially be duplicated.

2.5 Related Work

2.5.1 eQUIC Gateway

There have been previous publications on making QUIC more efficient by using BPF programs such as [PVV21], where a BPF program is used together with the Linux eXpress Data Path (XDP) to filter packets based on information provided by the userspace. This approach provided significant performance improvements with an increase of throughput by almost a third and a reduction of CPU time consumption caused by filtering packets by more than 25%. This shows that a setup leveraging Linux kernel features such as BPF has a lot of potential to improve the current infrastructure.

2.5.2 Kernel Bypass

Another interesting approach that follows a similar idea of speeding up packet processing by avoiding the Linux network stack is [Tyu22]. The difference in this work is that

DPDK is used to bypass the network stack to then process packets in userspace instead of using BPF programs like this thesis does. This, for example, offers more flexibility as the userspace program is not as limited (e.g. by the BPF verifier) as the eBPF program but might also lead to slightly more system calls, especially in the setup of a system, when user- and kernel-space need to communicate.

2.5.3 Priority drop

The idea of dropping packets based on their priority to adapt a connection in a congestion event has also been around for a while. [KWF03] explores this in more detail. Mainly improvements like a more tailorable congestion handling than the sole usage of discrete video quality levels as well as an improvement to, potentially randomized, frame dropping are discussed. This thesis, similar to [KWF03], will not focus on how the priority for packets is determined but rather on how those marked packets are handled. For this, it is assumed that a higher level protocol has correctly determined the packet priorities and can handle the drop of packets with lower priority in case of limited bandwidth.

3 Fast-Relays

In this chapter, we will dive into the specifics of our proposed fast-relay setup. We will cover necessary library adaptations as well as setup specifics like eBPF programs and userspace synchronization. Figure 3.1 already shows a high-level overview of a basic server-relay-client setup with its respective networking layers. A conventional setup would have the packet travel into userspace at the relay, while our setup aims to reduce the critical path, as indicated by the red arrows.




Figure 3.1: Conventional networking layers a packet passes. The red loop indicates again the “short-cut” that is utilized by the fast-relay (eBPF packet-forwarding - no need to go up to userspace).

3.1 QUIC Adaptions

As was already mentioned in the previous chapter, our setup requires some adaptations to the quic-go library. One initial change necessary was to turn off packet en- and decryption, that is happening within quic-go. Given that we operate on the QUIC-header data within the eBPF-program, we need access to fields that are encrypted using QUIC header-protection. For obvious reasons sending unencrypted packets is not something that would be wanted in a production environment, but for our exemplary

setup, it is unproblematic. An alternative would be to ‘push down’ en- and decryption onto a smartNIC using a hardware offload, but at the time of writing, there was no such offload available for QUIC. Given that such a hardware offload is added in the future, the en- and decryption can be turned on again, which makes this change more of a temporary solution to show the feasibility of our approach.

3.1.1 Function-Pointer Style Additions

Another type of change that we needed to introduce into the quic-go library is caused by connection state management. We essentially added support for communication with the eBPF-program by using an approach similar to C-style function pointers. On multiple locations, we added conditional function calls like the one depicted in Listing 3.1. The function that is called here will be defined by the developer of the relay and therefore allow for customizability without the need for changing the library itself.

```
1 /* Function pointer call within actual quic-go code */
2 if packet_setting.ConnectionUpdateBPFHandler != nil /* && potentially other
   conditions */ {
3     packet_setting.ConnectionUpdateBPFHandler(connId.Bytes(),
         uint8(connId.Len()), p.connection)
4 }
```

Listing 3.1: An example of a function-pointer addition to the quic-go library.

```
1 /* Function pointer signature definition within additional config file */
2 ConnectionUpdateBPFHandler func(id []byte, l uint8, conn QuicConnection) = nil
```

Listing 3.2: Only the signature will be defined within the library itself.

The definition of the function that the developer of the relay wishes to be executed at specifically defined points will be defined locally in the relay code and provided to the configuration of the quic-go library. An example of how this could look like is shown in Listing 3.3.

```
1 /* Definition of the function within the local relay code */
2 func localUpdateConnectionId(id []byte, l uint8, conn
   packet_setting.QuicConnection) {
3     /* handle the connection update by interacting with the eBPF-program */
4 }
5
6 /* Providing the function to the quic-go library */
7 func main() {
8     /* ... */
9     packet_setting.ConnectionUpdateBPFHandler = localUpdateConnectionId
10    /* ... */
11 }
```

Listing 3.3: An example of how the addition looks on the relay side.

The need for these additions arises since the eBPF-program works with its own copy of the current state of a connection. This, for example, includes the connection-id that will be used when changing the packet header before sending it out. Since a connection-id can change, i.e. be updated or retired, during the lifetime of a connection, we need a way to inform the eBPF-program to no longer use outdated state-information. These function-pointer style additions provide a minimal way of adding such functionality without limiting flexibility or adding too much application-specific code to the library itself, as would be the case if the library would access the eBPF-Maps directly.

3.1.2 Direct Changes to the Library

Besides the simple function-pointer style additions, we also had to make some direct changes to the library. These include simplifications of the packet structure to make the implementation of a prototype easier but also changes that are necessary for the whole approach to work. The necessary state changes mainly required internal state adjustments that would not be possible from outside the library because of missing access/visibility. The optional turn-off for the packet en- and decryption based on the value of the CRYPTO_TURNED_OFF flag also required some direct changes to the library, but these will not be mentioned further since they are temporary and not a focus of this chapter.

Simplifications of Packet Structure

We added some changes to the quic-go library to make a prototype implementation easier. The first one is that we fixed the sizes of some variable length fields like the connection-id or the stream-id. This was mainly to avoid the need to figure out the correct sizes within the eBPF-program as this would have resulted in approaches that would have to be very carefully turned into eBPF-compatible code due to verifier restrictions. Besides fixing the length of some fields, we also limit the number of

frames per packet to one. Normally a QUIC packet can contain multiple frames, especially stream-frames, but this would have required a packet traversal within the eBPF-program that, again, would have been harder to implement considering all verification constrictions. Besides this additional complexity, for our usecase the “single-stream-frame-per-packet” approach is not changing a lot since the payload of one media packet will be split into multiple packets due to its size. This means a single stream frame is likely already using up a whole packet, and multiple frames per packet would not be happening.

Stream Priorities

The first change, which is not only for simplification of our prototype implementation but is actually needed for our approach to work, is the addition of opening a stream with a specific priority. Our assumptions define that the server is the one marking the packets with the correct priority, which in our case means sending them via the correct stream. Such a stream is bound to a specific priority, and our code changes will make sure that the connection-id that is used for any packet sent on this stream will contain the correct priority value. Figure 3.2 visualizes the setup as well as our rudimentary approach of saving the priority value as the first byte of the connection-id.

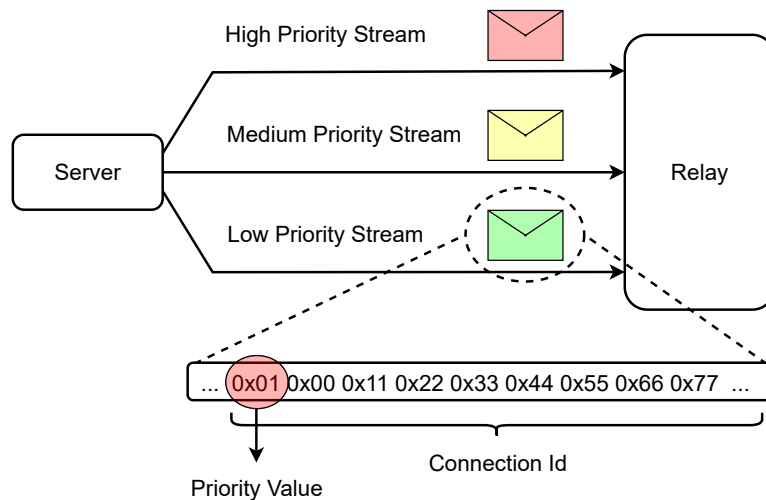


Figure 3.2: The server opens streams with specific priorities, which are then used to send packets with the corresponding priority.

This approach of having the priority value saved within the connection-id caused the need for another change in the library. Due to the periodic change/retirement of connection-ids, we need to make sure that at any point in time, there exists one

connection-id for each priority within the connection. Otherwise, we might be unable to mark a packet with the correct priority. This led us to introduce some additional logic to the code that is executing the retirement of connection-ids. There, we will make sure that before a connection-id is retired, either one with the same priority value already exists or is created. This solves the problem of not being able to mark a packet with the correct priority.

Retransmission

Another direct change that is needed is that of a specific `OnLost` function for packets that have been forwarded by the eBPF-program. Since the relay state is not necessarily aligned with the server state, and the relay does not know about the stream a packet was sent on, it is not possible to reuse the plain `OnLost` function of the quic-go library. Instead, our new function needs to look up the stream-id of the lost packet and open a new stream with that same id to ensure that the client correctly receives the retransmission. The relay also needs to tell the underlying eBPF-program that a packet is part of a retransmission and that, even if a stream was actually created by the relay, the egress program should treat it like it was created by the server. That last part is necessary for the stream-id translation part of the egress program, which is explained further in section 3.2.1, to work correctly. This functionality of informing the eBPF program about the retransmission, however, is again realized by a function-pointer style addition and thus designed by the relay developer.

Visible Endpoint for Packet Registration

The last direct change that we added was the introduction of an additional function `RegisterBPFPacket` on a quic-go connection object that allows the relay to register a packet. Since the packet registration requires access to internal state of the connection, this also needed to be done as an actual change to the library. Now the relay can just read the necessary information of a packet that needs to be registered from the eBPF-maps and then pass it on to this function which will handle the registration. This also provides a good separation between the Go code that handles eBPF communication and the actual QUIC connection handling.

3.2 eBPF Setup

3.2.1 Different eBPF Programs

In order to allow the relay to forward packets independently of the userspace, we need to equip the relay with three BPF programs, as seen in figure 3.3. Those three programs are:

- a program that handles incoming traffic **from** the clients (client ingress)
- a program that handles incoming traffic **from** the video server (server ingress)
- a program that handles outgoing traffic **to** the clients (client egress)

Their responsibilities then are:

- handling the initial registration of new clients and storing their information, such as MAC addresses in a BPF map
- intercepting the packets from the video server, duplicating and redirecting them to the egress program (as well as sending one unaltered packet to userspace for state management purposes)
- receiving the redirected packets at egress, altering them using the client-specific data, deciding (based on packet priority and client congestion) if a packet should be dropped or sent, storing info on sent-out packets for future congestion control purposes, and finally, sending them out to the clients

This setup allows us to separate any state management and congestion control from the actual packet forwarding and thus makes leaving out any immediate userspace processing possible.




Figure 3.3: The relay has to be equipped with three BPF programs.

Following is a more detailed description of the responsibilities of each of the three programs.

Client Ingress

This ingress eBPF program initially does some simple packet inspection on every incoming packet looking if the packet uses the correct protocols and addresses the right application layer. This is done by initially parsing the Ethernet, IP, and UDP headers, if present, and checking if the port matches the one the relay application is listening on. This means the correct port is to be defined prior such that the eBPF program can associate a single or (multiple) relay instance(s) with the correct port. In our case, then, since we use QUIC, the program will check for QUIC long header packets that setup

an initial connection and save the transmitted state information such as connection-id, stream-related states, etc. in an eBPF map together with information that is not directly known by the userspace, such as the client's MAC address. Saving data like the MAC address directly once the connection is set up allows to omit any further Address Resolution Protocol (ARP) steps later on.

Server Ingress

Another ingress-related program, this time for packets coming from the video server, is needed to handle packet duplication and forwarding. This program will receive the actual video packets from the server and then, based on an internal counter of how many clients actually want to receive the video, duplicate the packet accordingly. The counter indicating the number of clients will be updated by the userspace once a new connection is fully established and the client is ready to receive the actual video data. This might potentially cause some minuscule delay when updating the counter, but sending cached video data to the client for a brief moment when updating the counter could be a solution. Figure 3.4 shows the packet duplication and forwarding process high level for an example setting with three clients that want to receive the video data. In this ingress program, we need to consider a few things to ensure the correctness of our approach. These are:

1. The program can **only** forward packets that contain video data and must not forward any other packets that contain e.g. control data.
 - a) This is fairly easy to achieve by doing some header inspection of the QUIC header, which contains the packet type. Also, since the payload is generally sent using 0-RTT- (i.e. short-) headers, there is no need to consider long-header packets.
2. The program should pass an unaltered copy up to userspace to allow the QUIC library to gain knowledge of the packet and handle any state changes accordingly.
 - a) Generally speaking, this is not strictly necessary as one could just have a separate setup of registering packets that came from the server, but as this is not needed, it is considerably easier to just pass the packet up to userspace and let the library handle it normally. This does not impose any additional overhead as the forwarding of any duplicate packets happens independently, of course.
 - b) Any packet that has been identified as not part of our dataflow should, of course, be passed up to userspace without being forwarded to egress.

Duplicating and forwarding incoming packets that were identified as containing payload can be done using `bpf_clone_redirect(skb, egress_ifindex, 0)` which is a

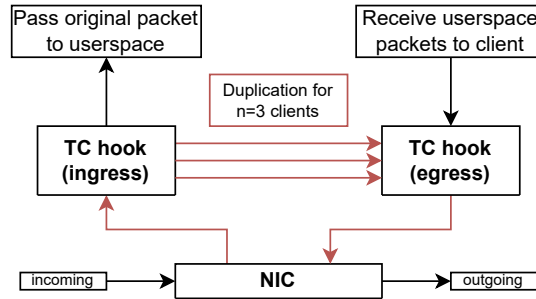


Figure 3.4: Duplication and forwarding of video packets to egress directly from ingress.

helper function that allows one to clone the packet buffer provided as the first argument and add it to the queue of the interface with the index that is provided as the second argument. The last argument allows for the specification of additional flags. Aside from `bpf_clone_redirect`, there are also other helper functions regarding packet redirection with slightly different behavior, so it is crucial to choose the appropriate one to get the desired outcome. Table 3.1 shows them together with a brief description of their behavior.

Helper	Description
<code>bpf_clone_redirect</code>	Clones and redirects a packet to the interface associated with the provided index.
<code>bpf_redirect</code>	Redirects a packet to the interface associated with the provided index. The packet is not cloned (no underlying call of <code>skb_clone()</code>) so it is slightly more efficient (25% pps increase according to commit message). The packet is also not redirected immediately but after the function finishes.
<code>bpf_redirect_peer</code>	Similar to <code>bpf_redirect</code> , but instead of redirecting the packet to the interface provided as a parameter, it is redirected to its peer device. This works only between different netns to allow for an efficient “ingress to ingress netns switch”. The switch is more performant since the packet does not need to go through the CPU backlog queue.
<code>bpf_redirect_neigh</code>	Again similar to <code>bpf_redirect</code> but allows to redirect a packet to another net device. This helper also fills in all the correct L2 addresses of the neighboring subsystem. Internally this executes a neighbor lookup to find the needed L2 information.

Table 3.1: Helper functions for packet redirection. There are more than that, but those are the ones we identified as possibly useful in the beginning.

Based on the descriptions of the redirection helper functions mentioned in table 3.1 it becomes clear that `bpf_clone_redirect` is the one suitable for our use case. This is because we need the cloning aspect since we essentially want to duplicate the incoming packets multiple times. Also, the redirection to another namespace or another net device, as provided by `bpf_redirect_peer` and `bpf_redirect_neigh`, respectively, is not needed in our case since we operate in the same relay-namespace throughout the whole process.

Client Egress

The central part of the eBPF setup where everything, from state-management to packet forwarding, comes together is the program that handles the outgoing traffic towards the clients. The client egress program sees every packet that leaves the relay, which includes packets that have been redirected by the ingress program as well as packets that have been generated by the relay itself (i.e. come from the relay userspace). This means that the program essentially merges two streams of packets into one stream that needs to be in a consistent state. This interleaving of packets grows the requirements of the program to the following list. Bold numbers indicate that a requirement is caused by the interleaving of packets.

1. Obviously, similar to the other eBPF programs, any packet that is not part of our traffic should be passed on normally without any further processing.
2. In case the packet is QUIC (for the correct client connection), the packet number needs to be changed to a program internal counter to avoid issues of reusing packet numbers. This is the only place where we can guarantee sequential packet numbers since neither the userspace nor the server knows what the highest packet number sent at any moment in time is. As there is no way of synchronizing lookups (e.g. using eBPF maps) of information between userspace and an eBPF program, changing it right before sending it out is the best way of keeping them coherent.
3. In case the packet (additionally to being QUIC for a client connection) contains stream frames, we need to do a similar translation of the stream id. The reasons and methods for this are the same as for the packet number translation. Figure 3.5 shows a flow diagram of the stream id translation process. Important steps include checking if the stream id translation already exists as well as checking if the packet is a retransmission. The former is necessary in case a payload is split into multiple packets, and the latter is necessary since the retransmission physically comes from the relay but should be treated as if it came from the server. For that the eBPF program will be notified by the userspace regarding which packets are retransmissions.
4. Again, if the packet is QUIC and for the right connection, the program needs to read the packet's priority and decide via a map-lookup if the connection state allows sending packets with the given priority.
5. In case the packet is a redirected media stream data packet, the program needs to find out which client the receiver is. This is done by saving the client id in some part of the packet that will be overwritten at egress (namely the connection-id) before redirecting the packet. The egress program, knowing where the id is saved in case of a redirected packet, can then lookup the correct address data of the

client (i.e. MAC address, IP address, etc.) using the saved id and overwrite the respective fields in the buffer before sending it out.

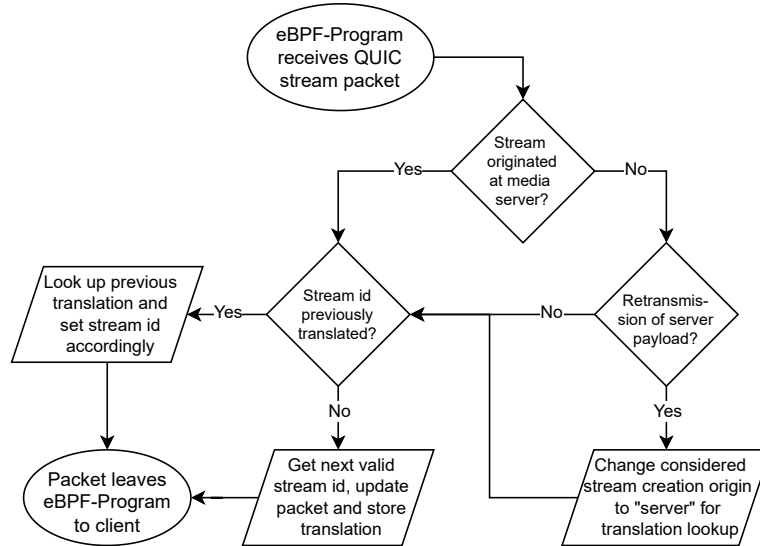


Figure 3.5: Flow diagram of stream id translation process.

The aforementioned packet number translation means that the packets sent by the userspace are likely to have a different packet number than the one chosen by the QUIC library. This might lead to inconsistencies again when receiving acknowledgments but can be avoided by remembering the translation in a map as well as only storing packet objects in the Go library that have the correct state. This technique of “storing” a packet that the userspace does not know about since it comes from the media server will be referred to as “packet registration”. This initially gives a brief window where a packet was sent out but is not saved in the history of the QUIC library. Once the packet is then processed by the userspace routine handling the registration, any incoming ACKs for this packet can be processed correctly. The next section will go into more detail on how the packet registration works. Later, we will also look at the implications of this on retransmissions. Even though the stream id translation works very similarly to the packet number translation, it does not have the need for any additional work after the packet has been sent out since our approach uses unidirectional streams only, and thus the relay does not care about changes in the used stream id.

3.2.2 Packet Registration

In order to make the congestion control algorithm that is running in userspace usable, we need to inform the QUIC library about the forwarded packets. This again happens

via BPF maps and a separate go routine that continuously polls new entries in the map and processes them. Entries are then added to the packet history to allow the receipt of ACKs. Besides that, the congestion control algorithm will be informed about the forwarded packet in order to be able to react to potential congestion events. Figure 3.6 visualizes the setup for this process.

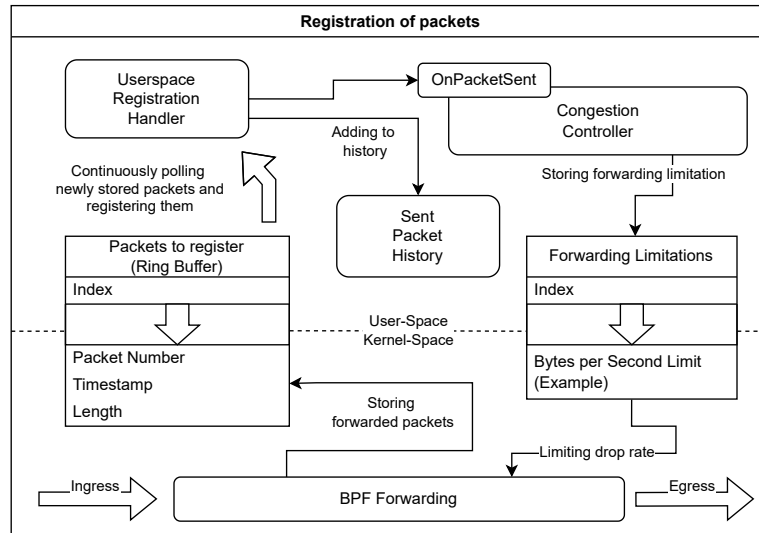


Figure 3.6: Internal setup for registering forwarded packets as well as incorporating forwarding limitations for the BPF program.

3.2.3 Retransmissions of Forwarded Packets

The fact that the relay only knows about the packets after a separate registration call causes another issue when it comes to retransmissions. Retransmissions happen at a stream level, which means that the relay needs to know about the stream in the first place if it wants to retransmit a packet. Given that we actually never create the streams that contain the frames of the payload data at the relay, this is a problem. Our proposed system handles this by manually creating a stream with the correct id the first time a packet goes missing for that stream. This can be imagined as a “just-in-time” stream creation before the retransmission is sent.

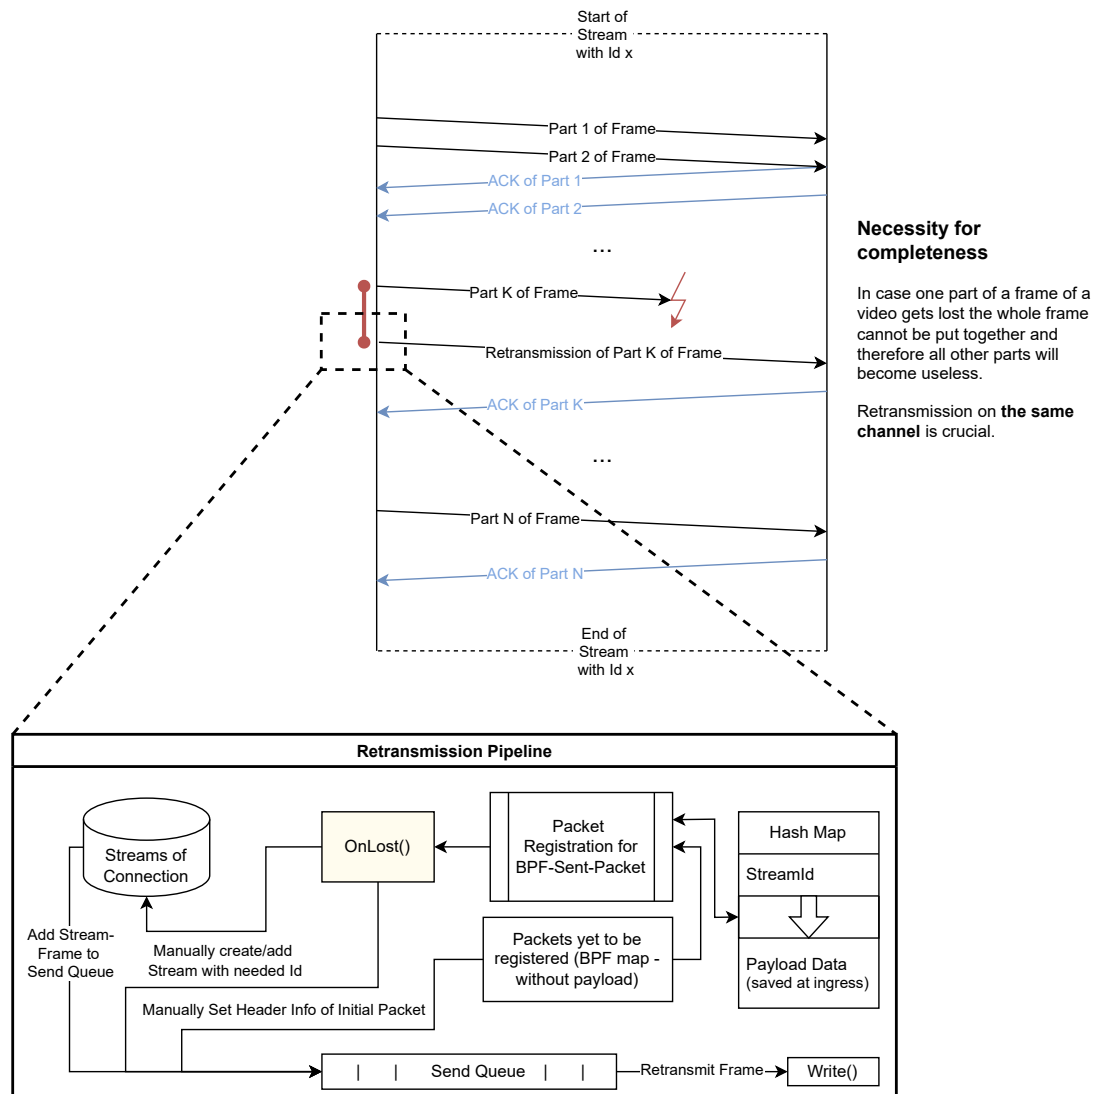


Figure 3.7: Internal setup for retransmitting packets forwarded from server connection.

It is not necessary to create the stream before that because our approach is only working with unidirectional streams so the relay does not get any data back from the client. The only case where it matters that the relay state “knows” about the stream is when a packet gets lost, and the `OnLost()` function of the QUIC library is called. This is also the latest point in time when the relay can create the stream with the needed stream id. Figure 3.7 visualizes the setup for retransmitting a packet that got lost on its way to the client.

3.3 Userspace Synchronization

Since one of the main ideas we propose is to avoid passing a packet all the way through the network stack up to userspace at the relay, we create the problem that the application itself is not aware of packets that are being sent. To conquer this issue, we suggest a setup that establishes communication between the application- and the eBPF-program. To keep the improvement we gained by not passing packets to userspace during the critical path, this communication will happen in a delayed fashion, that is decoupled from the actual sending of media data.

The main resource we use for this communication will again be eBPF maps that will contain information on the time a packet was sent together with protocol-specific data like packet numbers or stream identifiers (for both packet number and stream id the old and the new, i.e. translated, values will be stored).

3.3.1 Subscription and State Management

As already mentioned in section 3.2, an eBPF program handling incoming traffic from the client will save client connection information like MAC address, IP address, etc. in a map for later access. Also, an internal counter will give each client a unique identifier. With that, the only thing that happens in terms of communication between the application and the relay in case of a new subscription is that the application will update the “number of clients” counter that is accessed by the eBPF program and used for packet duplication purposes.

Regarding stream state management, there is also little need for communication since the server is expected to use QUIC’s unidirectional streams for sending the media data. That means the relay does not really need to know about the stream details except in case it has to trigger a retransmission. If that is the case, the stream id contained in the packet (meta-)data that was read from the eBPF map will be used to manually create a stream with the correct id. It is important to manually set the correct id since the relay might not use the same id for the next unidirectional stream it opens. Also, the client expects the retransmission to be sent on the same stream id as the original packet since a retransmission happens within the same stream-context.

3.3.2 Relay Caching

Regarding the caching of data within a relay, which is required by the MoQ standard, we can easily handle this since we are passing on an unaltered copy of any incoming packet from the server to the application at the same time we forward all the other packet copies to egress. Now the application is able to receive any data from the server as if this was a normal connection and store, e.g. the last n second(s) worth of data in a cache. Then the relay could, once a new connection is established, parallel to incrementing the kernel counter representing the number of clients also send out

cached data already so that the client receives it as early as possible.

One aspect that we still left open is the point in time when the relay should stop sending cached data since the forwarded data is up-to-date. This also includes the question of how the client can handle potentially duplicate media data if the cached- and the forwarded data happen to overlap. Such questions would likely require further experimenting and testing to find a good solution.

3.4 Congestion Considerations

QUIC, like many other modern transport protocols, contains congestion control mechanisms regulating the rate at which data is sent to a client. This is primarily done to avoid the network becoming congested, but it also has the secondary effect of circumventing the problem of overwhelming the client. Simply forwarding all packets the relay receives from the server would cause the relay-client connection to no longer have its own congestion control. Rather the rate at which the relay sends/forwards to the client would be determined by the server's congestion control algorithm, i.e. the network congestion between server and relay. Obviously, this is not a desirable situation, so our approach suggests the eBPF program at egress to have its own congestion control functionality.

3.4.1 Client Congestion

Already hinted at in figure 3.6, it is shown that once a packet is registered, there will also be a map update that will be triggered by the congestion controller. This map update will tell the eBPF egress program how much data it is allowed to send out. In figure 3.6 this is visualized exemplary as "Bytes per Second Limit" but the idea is that both the function determining how limits and thresholds are calculated from the information on incoming packets as well as the actual handling within the egress program will be application-specific and defined by the relay engineer. We experimented with approaches that use the QUIC internal measurements like the RTT, introduce new measurements like exponential-weighted moving averages, or even use an out-of-band connection where the relay expects direct feedback from the client. All these possibilities showed us that there is a lot of room for experimentation and optimization in this area. This, however, will not be explored further in this thesis and is left for future work.

3.4.2 Packet Filtering and Dropping

Assuming that the network congestion state is known and communicated to the relay, one can use the priority-information within a packet (that is expected to be set by the server) to decide which packets should be forwarded and which ones should be

dropped. One difficulty in this approach is that the dropping mechanism essentially works as an online-algorithm, meaning that it does not have full knowledge of the traffic, especially not of future packets. This means that a case like the following could happen:

- The traffic contains packets within the priority range of 1 to 5 (5 being the highest priority).
- Given the current network congestion to the client, the relay decides to drop all packets below priority 3 and 50% of the packets with priority 3.
- The remaining byte limit to be sent out is running low, and a packet with priority 3 comes in, which is sent, since the relay already dropped a lot of previous priority 3 packets.
- The next packet turns out to be a priority 5 packet which would overshoot the byte limit if sent.

In this example, one could use many different heuristics to handle the situation. One could always keep enough byte limit left so that a high-priority packet can always be sent. This, however, essentially just lowers the byte limit for all other packets while making it higher for high-priority packets. Therefore, an individual byte limit per priority could also be used right away. Also, this approach might cause problems if a lot of high-priority packets come in at once, e.g. in case of very “bursty” traffic. Another way that could be used to handle this uncertain situation is to allow for temporary overflows of the byte limit. This would make the limit more of a soft limit that can be exceeded for a period of short time. However, this is another heuristic highly dependent on the specific use case and traffic patterns, so we did not implement it in our prototype. Overall we can say that implementation-wise, it is not hard to drop packets, but the actual difficulty lies in finding a reasonable way of deciding which packets to drop.

3.5 Integration and Prototype

To show the feasibility of our approach, we have built a prototype that implements the suggested setup. It is capable of streaming an example video as well as mock-payload data from a server to a client via a relay. This setup will be open source and available on GitHub. Some of the mentioned functionalities, like retransmission, caching, and adaptive bitrate streaming, have not yet been fully implemented but should not require major additions since the infrastructure already considers them. The absence of these functionalities will not have a drastic impact on further testing, however since the retransmission and the adaptive bitrate streaming will not be needed in our lab-environment, and the caching is not expected to have a big impact either.

3.5.1 Compatibility

The application layer of the prototype is written in Go, but this is not a requirement itself. The eBPF program expects QUIC, but this could also be changed without too much effort to support any other underlying transport protocol.

3.5.2 Source Code Repositories

For the development of the relay and the eBPF programs, we have come up with the following repositories:

- **Fast-Relay** [Pfe24b]: This is the main repository providing the eBPF program implementations as well as examples of server, relay, and client implementations in Go.
- **Quic-Go Adaptation** [Pfe24a]: This repository is a fork of the QUIC library “quic-go” [See24] and provides a plain Go implementation of the QUIC protocol. For our thesis, we needed to make some adaptations to the library to support some hook points for separate functions, which should be specifically designed to handle the underlying eBPF setup with its eBPF-Map usage.
- **MoQ-Transport Adaptation** [PE24]: This repository is a fork of the “MoQ-Transport” [Int24a] protocol repository and provides some needed adaptations to our examples. One such adaptation is that the server needs to support a categorization of payloads into different priorities in order for the eBPF program to be able to deliberately drop packets in case of congestion. Getting these priorities could be as simple as differentiating only between I- and P-frames in a video stream or more complex based on the needs of the application and the wanted granularity of the congestion control.

4 Testing

4.1 Setups

We test the performance of our prototype in a lab-like environment that allows us to maintain a stable network state. We do this by using Linux network namespaces where each namespace represents a different participant in the communication, i.e. the server, the relay and the client. A schematic overview of the setup is shown in Figure 4.1.




Figure 4.1: Logical setup including interfaces and IPs

4.1.1 Local Environment for Testing and Development

The local environment we use for testing and developing allows us to setup the network in a way that fits the current needs of the prototype. One example for such needs would be that we need a slight delay between sending a packet and receiving its acknowledgment since the registering of a sent packet takes place after a packet is sent out. In a real-world scenario this would obviously also be the case, if the distances covered are large enough.

However, where a local setup turns out better for testing is its deterministic nature. The fact that the delay of a bridge can be set to a fixed value allows us to better analyse and understand the impact of the different setups for example on packet jitter. Also we can be sure that delays, and potential improvements thereof, are not caused by external network state but rather by the immediate forwarding our approach introduces.

4.1.2 Physical Server Setup for Real-World Testing

TODO: mention that the physical server posed problems because of an old kernel?




Figure 4.2: Local setup including bridges and namespaces

4.2 Testing and Results

When testing the performance of our prototype we will focus on showing that the eBPF forwarding is capable of reducing the delay of packets. Delay in our case is measured again using eBPF programs that save the timestamps of a packet leaving and entering the server and client program respectively. Due to the highly deterministic nature of the local environment we know that the difference between the two timestamps will be mainly due to processing at the relay. Besides delay analysis we will also show that our approach does not require more CPU resources than the plain userspace forwarding. For that we will look at profiling data from programs like “pprof” as well as the Linux “process status” (ps) command which can be used to show the CPU usage of a process.

4.2.1 Delay Reduction of eBPF Forwarding

When looking at the impact of eBPF-Forwarding on the delays of packets, we can see that in the namespace environment the delay of a single packet is decreased by around

100 μ s when compared to the userspace forwarding. This is shown in Figure 4.3 and considers the simplest case of the userspace program where it only forwards directly to one connection without the need for much additional computation. Given that the userspace can have arbitrary complex connection management this delay improvement can likely be even bigger in a more sophisticated userspace setup.

Another thing the figure shows is that the delay has a smaller variance due to the fact that the eBPF program path is somewhat similar for each packet whereas, in contrast, the userspace path can have buffers, queues, or similar that lead to a higher difference in processing time between packets. This effect however might be less observable in a real world scenario due to the ubiquitous network jitter which was influential in the used namespace environment. TODO

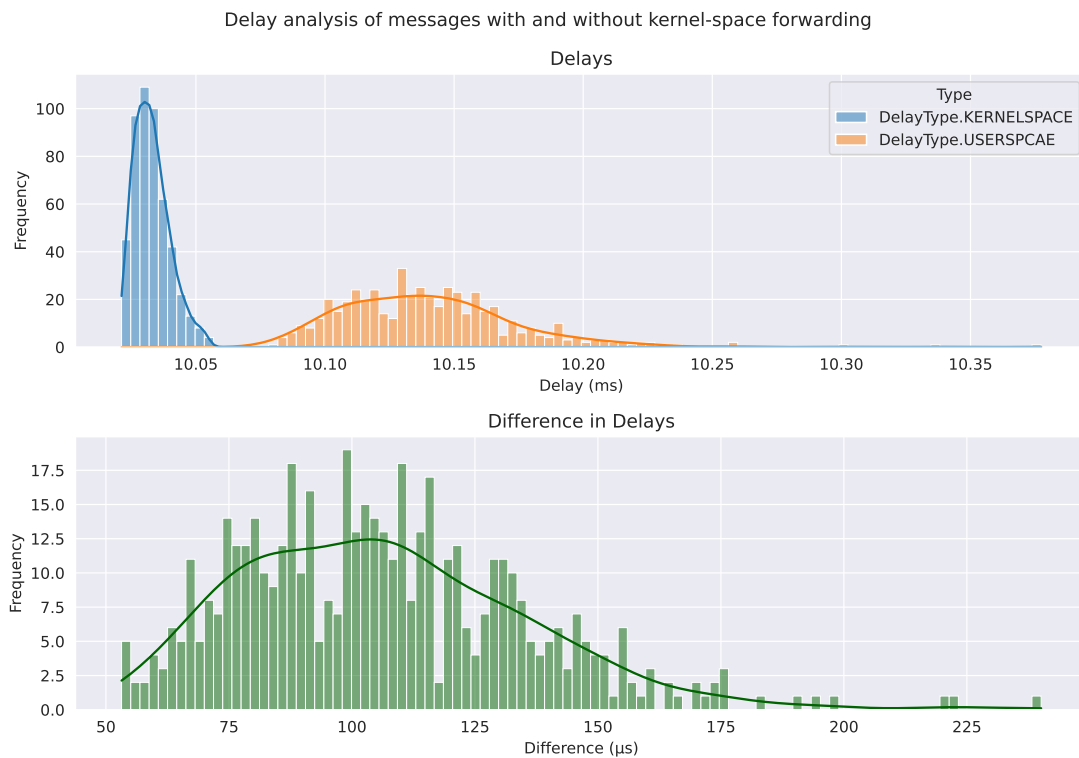


Figure 4.3: By avoiding userspace when processing the 1-RTT packets that contain the payload the delay of a single packet can be reduced. The longer-delay-operations are either handled directly in the eBPF program (e.g. the case for deciding where to redirect a packet to) or handled after the packet has already been sent out (e.g. the case for registering a packet such that the QUIC library knows about it).

4.2.2 CPU Utilization Comparison

TODO




Figure 4.4: Server namespace CPU usage comparison.




Figure 4.5: Relay namespace CPU usage comparison.




Figure 4.6: Client namespace CPU usage comparison.

TODO

Table 4.1: Table for CPU usage of relay Go processes

flat	flat%	sum%	cum	cum%	cause
440ms	28.39%	28.39%	440ms	28.39%	runtime/internal/syscall.Syscall6
140ms	9.03%	37.42%	140ms	9.03%	runtime.futex
120ms	7.74%	45.16%	120ms	7.74%	runtime.cgocall
40ms	2.58%	47.74%	40ms	2.58%	runtime.write1
30ms	1.94%	49.68%	60ms	3.87%	runtime.checkTimers
30ms	1.94%	51.61%	50ms	3.23%	runtime.mapaccess2
20ms	1.29%	52.90%	70ms	4.52%	.../cilium/ebpf/internal/unix.Syscall
20ms	1.29%	54.19%	20ms	1.29%	net.IPString
20ms	1.29%	55.48%	20ms	1.29%	runtime.casgstatus
20ms	1.29%	56.77%	20ms	1.29%	runtime.duffcopy

5 Future Work

5.1 Hardware Offload

This thesis heavily relies on the fact that the relay can access certain fields (e.g. packet numbers) of the packet, which are generally not accessible without prior decryption. In the current setup, this is made possible by turning off encryption altogether but to be of any use in a real-world scenario, the encryption of incoming and the decryption of outgoing packets would need to be pushed down below the lowest used BPF hook point in the stack. This means that a hardware offload of encryption and decryption similar to what is done for TCP/IP checksums would be necessary.

Once compatible SmartNIC offload implementations are available one can, besides en- and decryption, also offload the BPF program itself. This then would provide another way of accelerating performance.

Some previous work in this direction has already been done since at least 2019 [Yan+20] but for the purpose of this thesis we did not find any suitable open-source implementation that would allow incorporation into our fast-relay example implementation.

5.2 Compatibility Expansion

This thesis used the QUIC protocol together with media over QUIC (MoQ) to demonstrate how fast-relays that circumvent userspace by utilizing BPF programs can be designed. However, generally speaking the design of fast-relays is not limited to any of these protocols and could be expanded given modifications to necessary fields are possible (i.e. not prevented by encryption) and there is a way to encode the priority of a packet in the packet itself. The latter point could always be realized by using part of the payload which forces a deeper packet inspection within the BPF program but avoids the need to fit the priority into the header of an existing protocol.

6 Conclusion

TODO

List of Figures

2.1	Hook points within network stack	6
2.2	Exemplary eBPF map definitions	7
2.3	Rough MoQ relay architecture	9
2.4	Non-caching relay traffic	9
2.5	MoQ relay traffic	9
2.6	Adaptive streaming schematic	10
3.1	Packet path schematic regarding network stack	13
3.2	Streams with specific priorities	16
3.3	Types of eBPF programs at relay	18
3.4	Video packet duplication	20
3.5	Stream id translation schematic	23
3.6	Packet registration schematic	24
3.7	Packet retransmission schematic	25
4.1	Logical setup including interfaces and IPs	30
4.2	Local setup including bridges and namespaces	31
4.3	Delay analysis of eBPF approach	32
4.4	Server CPU usage comparison	33
4.5	Relay CPU usage comparison	33
4.6	Client CPU usage comparison	33

List of Tables

2.1	Subset of eBPF map types	8
3.1	Redirection helpers for packet buffer	21
4.1	Table for CPU usage of relay Go processes	34

Bibliography

- [Fou] L. Foundation. *eBPF Verifier*. Accessed: July 13th, 2024. URL: <https://docs.kernel.org/bpf/verifier.html>.
- [Int23] Internet-Engineering-Task-Force. *Media over QUIC (moq)*. Accessed: July 18th, 2024. 2023. URL: <https://datatracker.ietf.org/wg/moq/about/>.
- [Int24a] Internet-Engineering-Task-Force. *Media over QUIC Transport*. Accessed: July 19th, 2024. 2024. URL: <https://datatracker.ietf.org/doc/draft-ietf-moq-transport/>.
- [Int24b] Internet-Engineering-Task-Force. *QUIC: A UDP-Based Multiplexed and Secure Transport*. Accessed: May 26th, 2024. 2024. URL: <https://datatracker.ietf.org/doc/html/rfc9000>.
- [Int24c] Internet-Engineering-Task-Force. *What’s the deal with Media Over QUIC?* Accessed: July 18th, 2024. Jan. 2024. URL: <https://www.ietf.org/blog/moq-overview/>.
- [JC20] M. Joras and Y. Chi. *How Facebook is bringing QUIC to billions*. Accessed: May 26th, 2024. Oct. 2020. URL: <https://engineering.fb.com/2020/10/21/networking-traffic/how-facebook-is-bringing-quic-to-billions/>.
- [KWF03] C. Krasic, J. Walpole, and W.-c. Feng. *Quality-Adaptive Media Streaming by Priority Drop*. June 2003. doi: 10.1145/776322.776341.
- [Lan+17] A. Langley, A. Riddoch, A. Wilk, A. Vicente, C. Krasic, D. Zhang, F. Yang, F. Kouranov, I. Swett, J. Iyengar, J. Bailey, J. Dorfman, J. Roskind, J. Kulik, P. Westin, R. Tenneti, R. Shade, R. Hamilton, V. Vasiliev, W.-T. Chang, and Z. Shi. “The QUIC Transport Protocol: Design and Internet-Scale Deployment.” In: *SIGCOMM ’17: Proceedings of the Conference of the ACM Special Interest Group on Data Communication* (Aug. 2017), pp. 183–196. doi: 10.1145/3098822.3098842.
- [Lie18] E. Liebetrau. *How Google’s QUIC Protocol Impacts Network Security and Reporting*. Accessed: May 26th, 2024. June 2018. URL: <https://www.fastvue.co/fastvue/blog/googles-quic-protocols-security-and-reporting-implications/>.
- [PE24] D. Pfeifer and M. Engelbart. *Media over QUIC Transport*. Accessed: May 26th, 2024. 2024. URL: <https://github.com/danielpfeifer02/priority-moqtransport.git>.

- [Pfe24a] D. Pfeifer. *A QUIC implementation in pure Go*. Accessed: May 26th, 2024. 2024. URL: <https://github.com/danielpfeifer02/quic-go-prio-packs.git>.
- [Pfe24b] D. Pfeifer. *Adaptive Media over QUIC Transport*. Accessed: May 26th, 2024. 2024. URL: https://github.com/danielpfeifer02/Adaptive_MoQ.git.
- [PVV21] G. Pantuza, M. A. M. Vieira, and L. F. M. Vieira. *eQUIC Gateway: Maximizing QUIC Throughput using a Gateway Service based on eBPF + XDP*. 2021. DOI: 10.1109/ISCC53001.2021.9631262.
- [Rod24] L. Rodriguez. *The Future of the Internet is here: QUIC Protocol and HTTP/3*. Accessed: May 26th, 2024. Jan. 2024. URL: <https://medium.com/@luisrodri/the-future-of-the-internet-is-here-quic-protocol-and-http-3-d7061adf424f>.
- [See24] M. Seemann. *A QUIC implementation in pure Go*. Accessed: May 26th, 2024. 2024. URL: <https://github.com/quic-go/quic-go.git>.
- [Sta94] I. O. for Standardization. *ISO/IEC 7498-1:1994*. Accessed: July 3rd, 2024. 1994. URL: <https://www.iso.org/obp/ui/en/#iso:std:iso-iec:7498:-1:ed-1:v2:en>.
- [Tig] Tigera. *eBPF Explained: Use Cases, Concepts, and Architecture*. Accessed: July 19th, 2024. URL: <https://www.tigera.io/learn/guides/ebpf/>.
- [Tyu22] N. Tyunyayev. *Improving the performance of picoquic by bypassing the Linux Kernel with DPDK*. Master's Thesis. 2022.
- [W3T24] W3Techs.com. *Usage statistics of QUIC for websites*. Accessed: May 26th, 2024. May 2024. URL: <https://w3techs.com/technologies/details/ce-quic>.
- [Yan+20] X. Yang, L. Eggert, J. Ott, S. Uhlig, Z. Sun, and G. Antichi. "Making QUIC Quicker With NIC Offload." In: *EPIQ '20: Proceedings of the Workshop on the Evolution, Performance, and Interoperability of QUIC* (Aug. 2020), pp. 21–27. DOI: 10.1145/3405796.3405827.